

# CS 6290

## Privacy-enhancing Technologies


Department of Computer Science

Slides credit in part from E. Ben-Sasson, D. Dziembowski, I. Eyal, R. Geambasu, and F. Greenspan, A. Judmayer, A. Juels, A. Miller, J. Poon, V. Shmatikov, D. Song, and F. Zhang


# Lecture 3 – Energy-efficient Consensus Mechanisms

Prof. Cong WANG  
CS Department  
City University of Hong Kong

# Recall: PoW Mining




*Pr [Success]* proportional to computing power



# Bitcoin PoW: Criticism

- Relies on an ongoing computational race
  - Honest majority of computing power is critical
- Burns energy (**Proof of Work = Proof of Waste?**)
  - 10+ GW (<http://realtimebitcoin.info>)
    - Ireland consumes 3.1 GW and Austria 8.2 GW



# Challenge: Replace PoW with Alternate Resource Lottery

- Other physical resources, with different properties?
  - Disk space
  - Useful computation/storage
  - Etc ...
- What about the coin itself?
  - “Virtual resource mining” --> Proof of Stake (PoS)

# Challenge: Replace PoW with Alternate Resource Lottery

- Other physical resources, with different properties?
  - Disk space
  - Useful computation/storage
  - Etc ...
- What about the coin itself?
  - “Virtual resource mining” --> Proof of Stake (PoS)

*First question: can we recycle the waste of Bitcoin and do something useful?*


# Greening Bitcoin

*Proof-of-useful-works*

# Repurposing Bitcoin Work

- Permcoin:

- Idea: a puzzle where hardware investment is useful, even if the work is wasted?




- Side effect:

- Massively distributed, replicated storage system



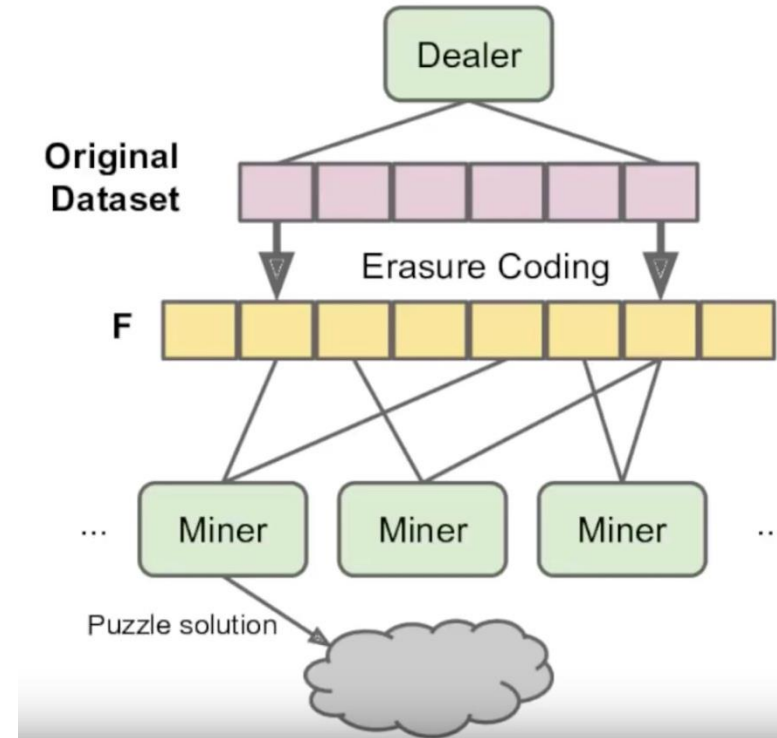
# Permacoin: Setting

- Assume we have a large public file  $F$  to store
- For simplicity:  $F$  is chosen globally, at the beginning, by a trusted dealer
- Each user (miner) stores a random subset of the file

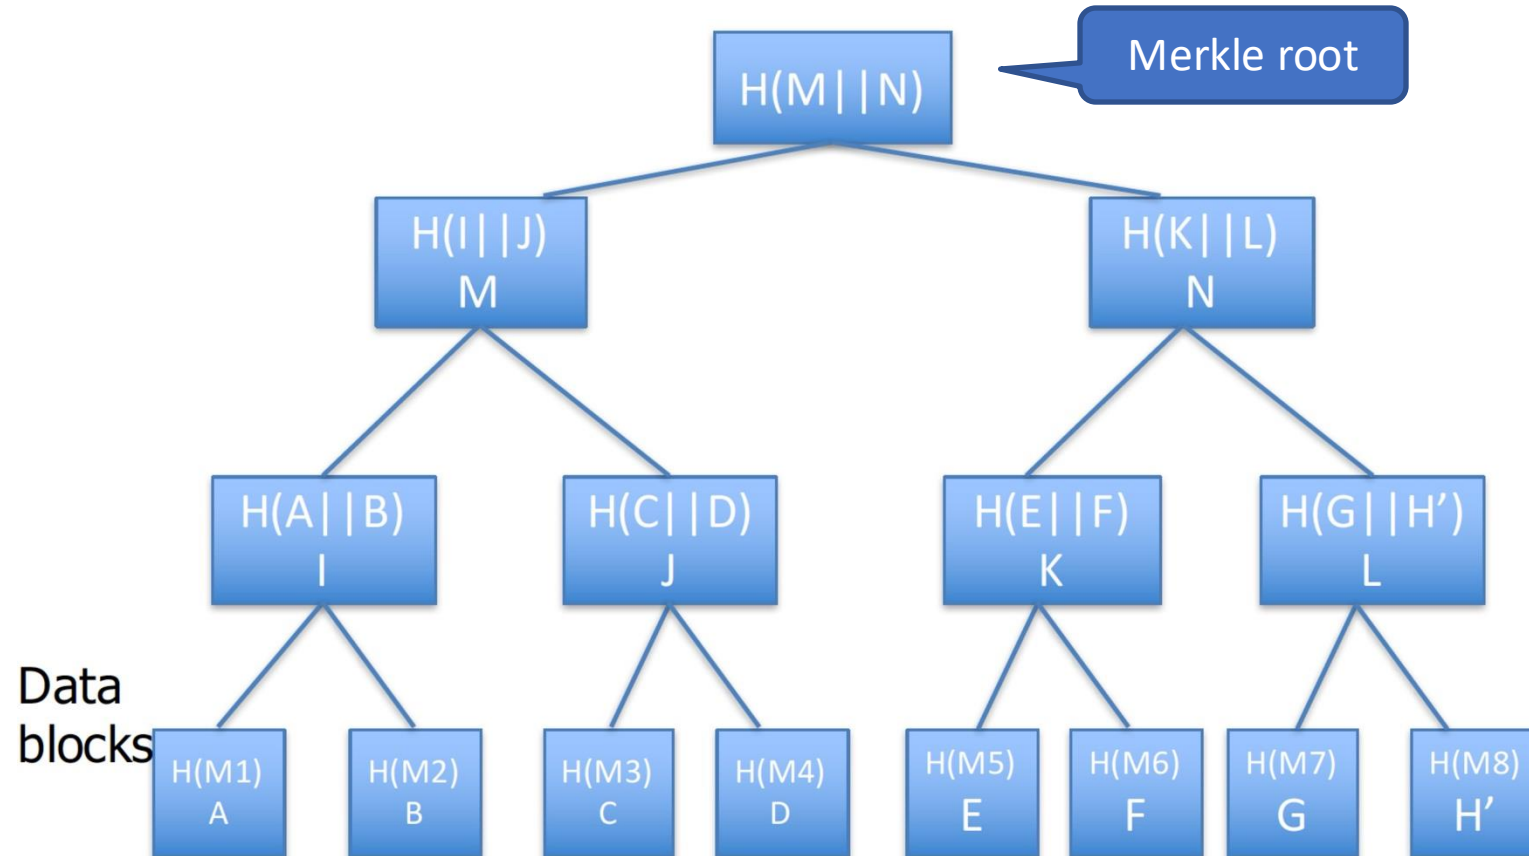


# Permacoin: Features

- **Storage:** a large public dataset
  - Optionally, users can submit their own data to archive
- **Recoverability:** even after catastrophic failure
  - Thanks to Erasure Coding (will introduce later)
- **Diversity:** geographical, as well as administrative
  - Store at individual miners

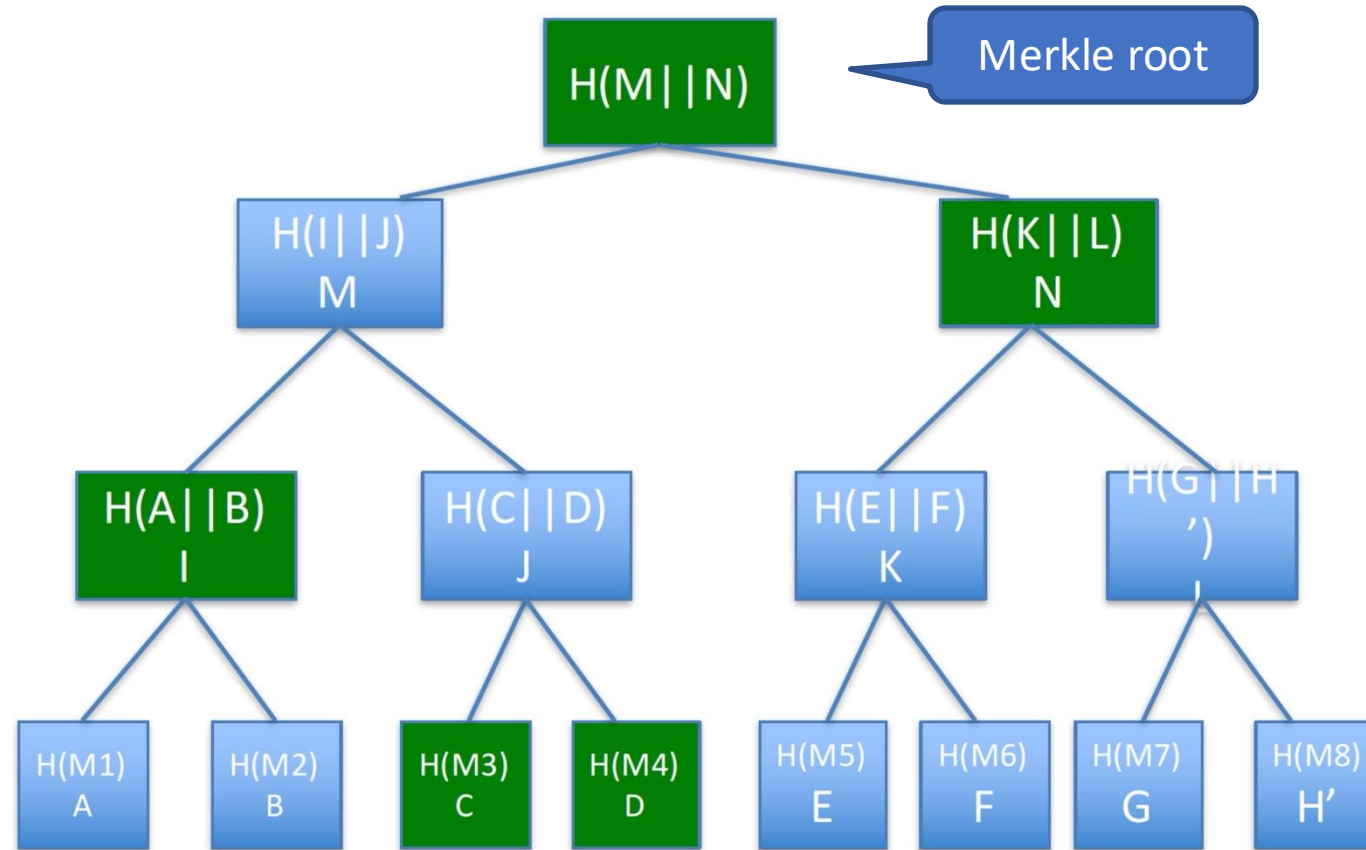


# Background: Merkle Tree



To check D, Proof =  $\langle H(M4), H(M3), H(A||B), H(K||L), H(M||N) \rangle$   
should match with root


# Background: Merkle Tree



To check D, Proof =  $\langle H(M4), H(M3), H(A||B), H(K||L), H(M||N) \rangle$   
should match with root

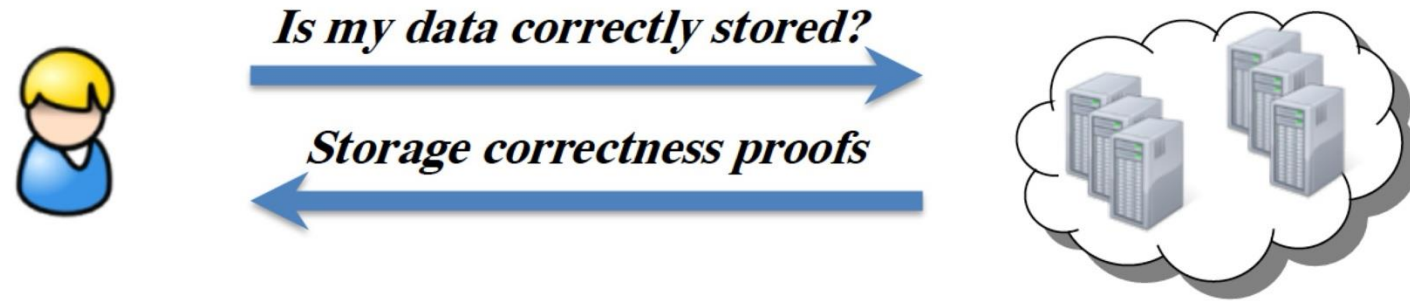
# Background: Erasure Codes

- $k$  data blocks  $\xrightarrow{\text{encode}}$   $k + m$  blocks
  - Example: Reed-Solomon 6+3



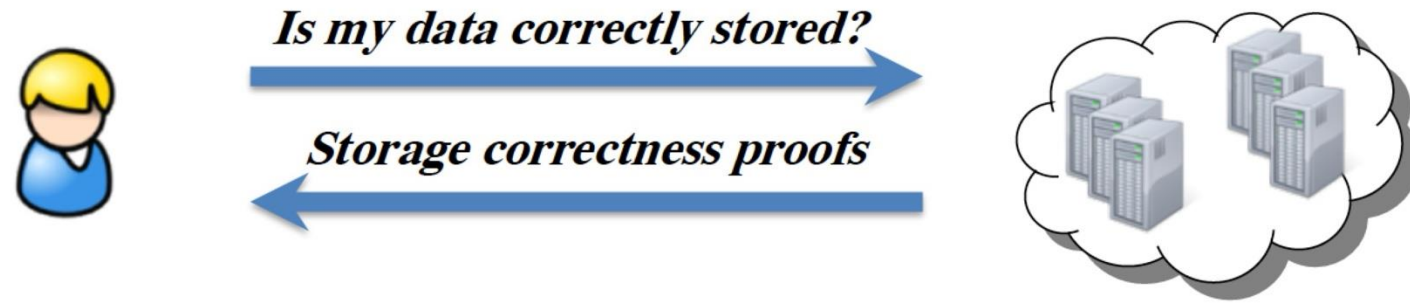
- Reliability: can tolerate  $m$  failures
- Efficiency:
  - Save disk space
    - Compared to previous effort: one data block replicate to many (say 3) replicas
  - Save I/O bandwidth on the write path

# Background: Data Auditing



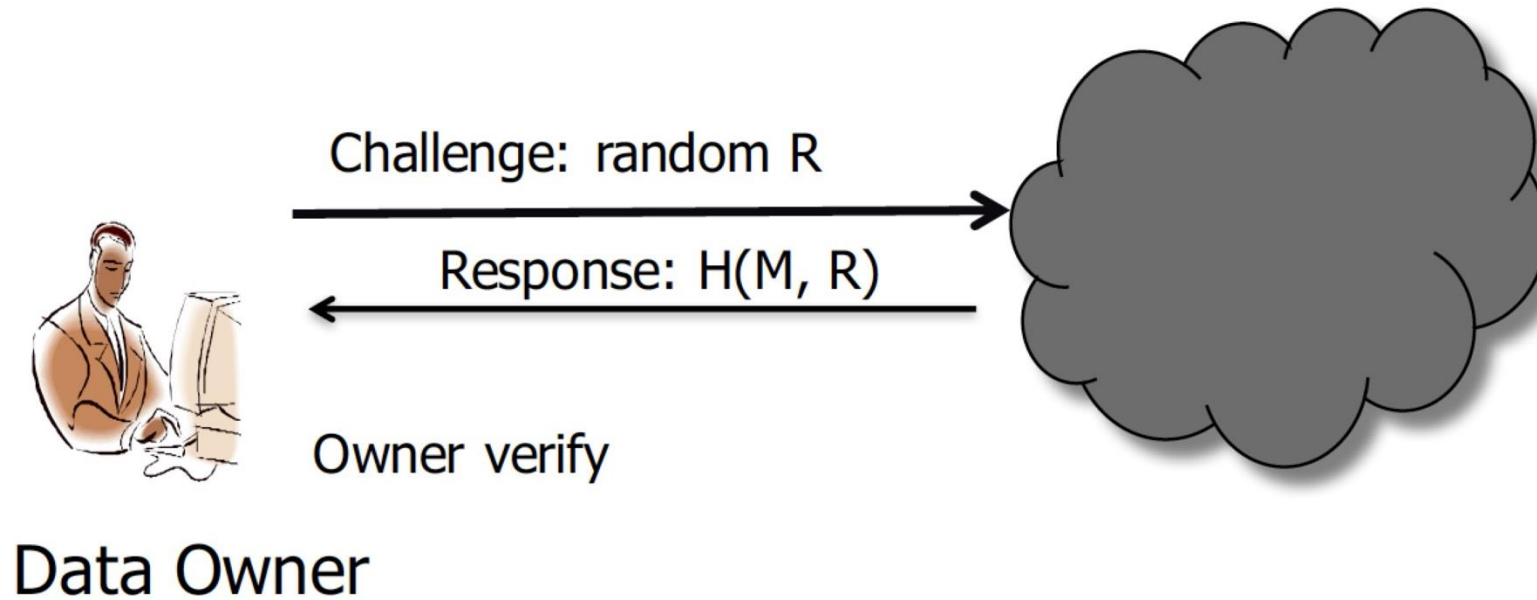
- Client stores data on the storage providers
- Providers can behave unfaithfully
  - Discard old data
  - Hide data loss
- Client needs a guarantee that data is stored correctly
  - To help extend data trust perimeter into the storage providers
  - To meet security, system, and performance requirements

# Background: PoR



- Proofs of Retrievability (PoR):
  - The client (verifier) runs an efficient data **audit proof** in which the data storage server (prover) proves that it **still possesses the client's data** and client can recover entire file

# How to Audit





# Secure Storage Auditing

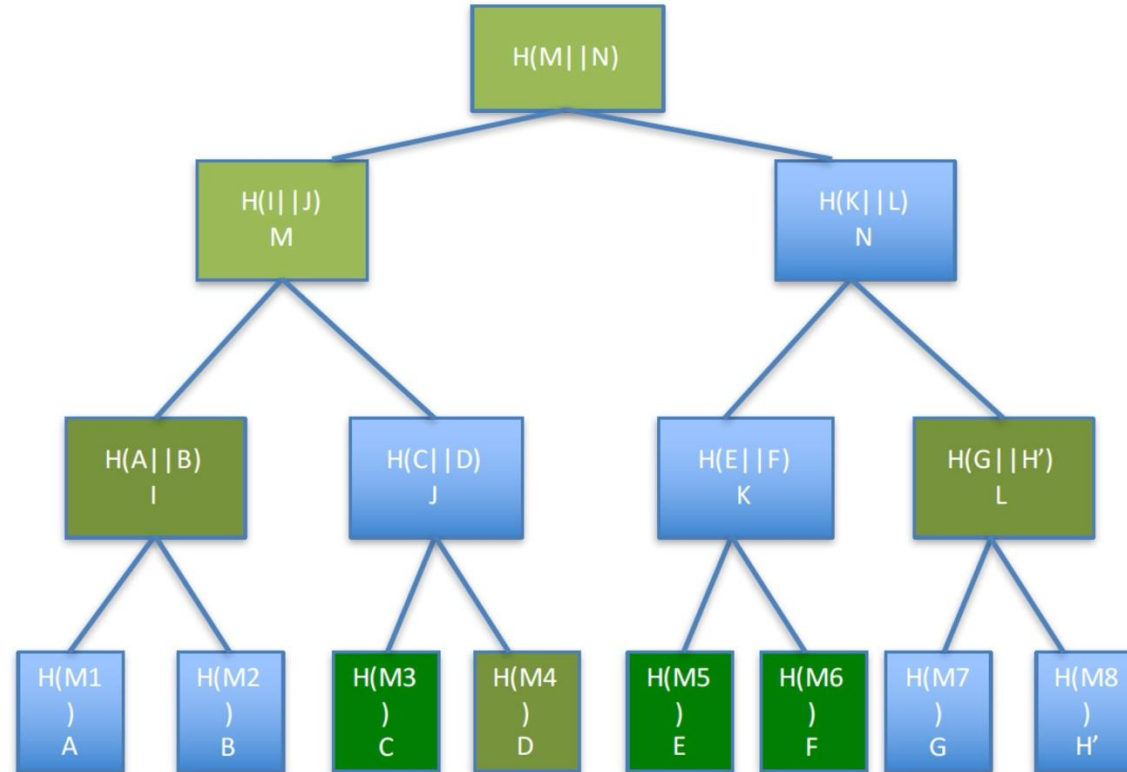
- Demand efficient storage correctness guarantee **without requiring local data copies**
  - Traditional methods for storage security can not be directly adopted
  - Retrieving massive data for checking is unpractical (large bandwidth)
- Allow meaningful tradeoffs between security and overhead
  - Communication and computation **costs** should be low
  - Auditing cost should not outweigh its benefits
- Cope with frequent **data changing** while ensuring continuous data auditing
  - Data may be frequently updated by owner for application purposes
  - Auditing mechanisms inherently need to support data dynamics

# Secure Storage Auditing (Cont.)

- Enable **public auditing** for unified risk evaluation
  - Introduce a third-party auditor (or miners in the Blockchain network) assists in data auditing
  - Public auditing should not affect owner's data privacy (**optional**)
- Handle multiple auditing tasks simultaneously (batch auditing)
  - The individual auditing of each data file can be tedious and inefficient
  - Batch auditing improves **efficiency** and saves computation overhead



# One Example: Merkle Tree for Data Auditing

- Challenge  
blocks: C, E, F
- Proof: D, I, L, M  
Root
- Verify with  
stored root  
value




# Data Auditing in Permacoin (1)

- Merkle-tree based
- The Merkle root **digest** of the file **F** is publicly accessible
  - Each leaf (e.g.,  $F_0$ ) is a segment of the file
- Each miner (with public signing key **PK**) is mapped to a random subset of segments to store
  - Basing on the hash of **PK**
  - E.g., **PK** here maps to four segments
    - $F_1, F_2, F_4, F_5$




# Data Auditing in Permacoin (2)

- Non-interactive challenge generation
  - Given the epoch-dependent puzzle **puz**
    - **puz** =  $v \parallel B_l \parallel MR(x) \parallel T$ , where  $v$  is the software version number,  $B_l$  is the previous mined block header,  $MR(x)$  is the merkle-tree root over new transactions  $x$ , and  $T$  is the timestamp
  - The random challenge indices are selected
    - a.  $h1 := H(\text{puz} \parallel \text{PK} \parallel s)$ , where  $s$  is a random string chosen by the miner **PK**
    - b.  $h1$  selects  $k$  segments from subset
      - E.g.,  $F_2, F_4$  are selected



# Data Auditing in Permacoin (3)

- The **ticket** for the puzzle **puz** is
  - $\text{ticket} := (\text{PK}, s, \{\mathbf{F}_2, \pi_2\}, \{\mathbf{F}_4, \pi_4\})$   
where  $\pi_i$  is the Merkle proof for segment  $F_i$
- **Last step:**
  - Compute  $h2 := H(\text{puz} \parallel \text{ticket})$
  - Winner if  $h2 < \text{TARGET}$



# Potential Benefits

- Reducing Bitcoin's “honest” cost
- One example: UTXO (unspent output from bitcoin transactions) database
  - Miners in Bitcoin validate every transaction
  - Validation requires the UTXO database (GBs)
  - However, currently maintaining the UTXO database doesn't pay
- **Idea: use Permacoin to reward UTXO storage**

*Besides augmenting the storage of a public dataset, it seems that we can do more ...*



# REM: Resource Efficient Mining for Blockchains

Fan Zhang, Ittay Eyal, Robert Escriva,  
Ari Juels, Robbert van Renesse



Vancouver, Canada

13 September 2017

USENIX Security 2017

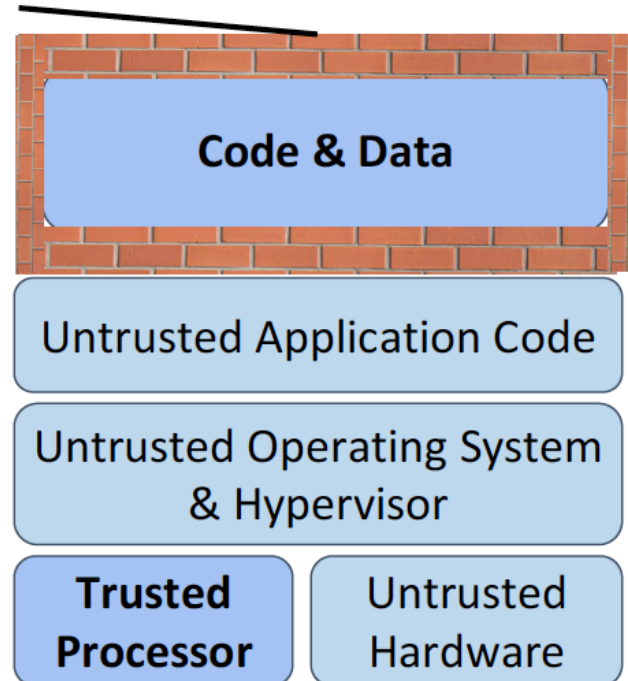
# Software Guard eXtension (SGX)

## Integrity



Other software and **even OS** cannot tamper with control flow.

“Enclave”

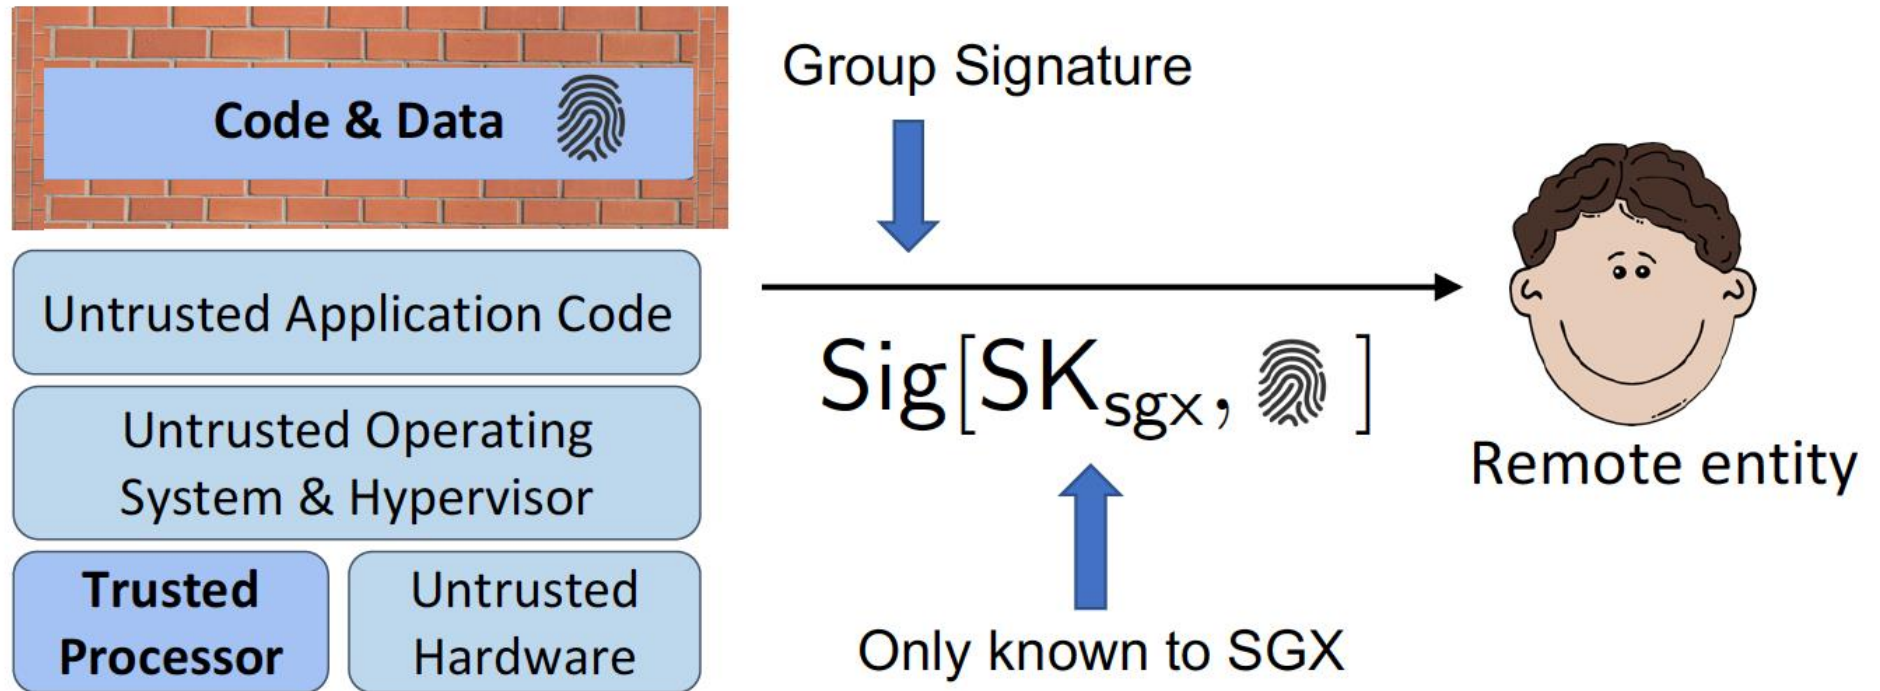


## Confidentiality



Other software and **even OS** can learn nothing about the internal state\*.

# SGX: remote attestation



# SGX-backed Blockchain: A New Security Model

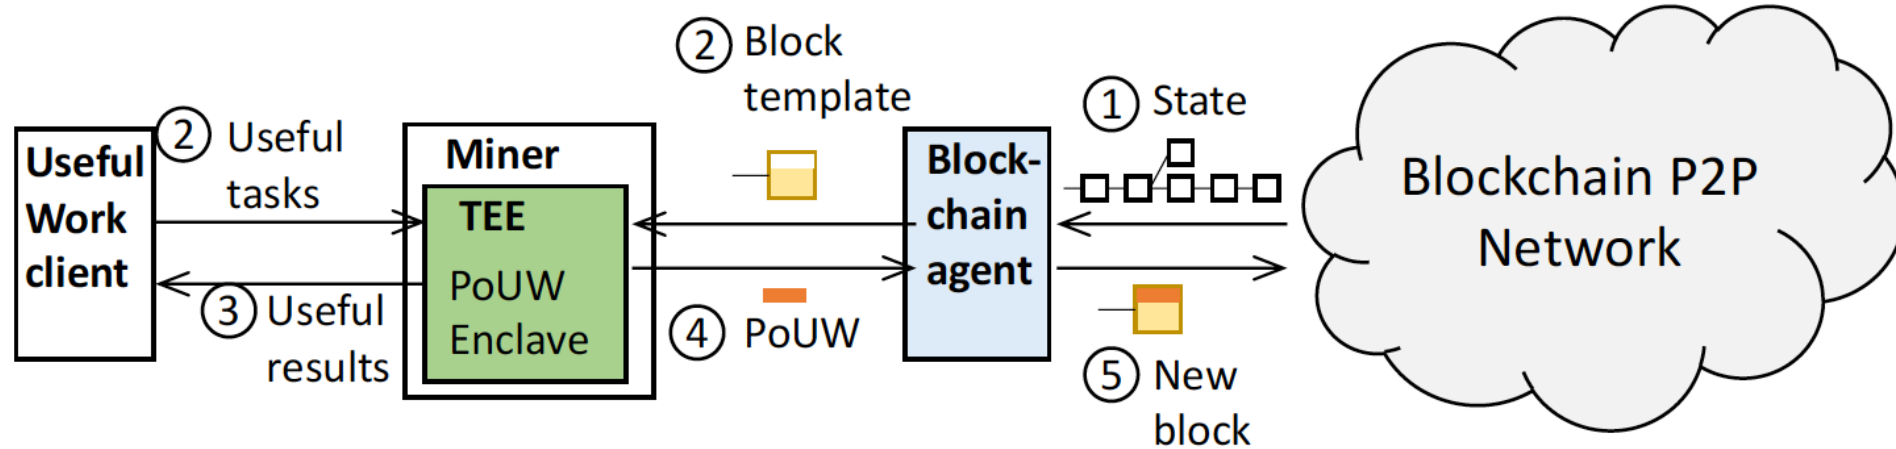
- **Permission-less**
  - Anyone with SGX equipped can join
- **Partially decentralized**
  - Assume SGX works as advertised
  - Intel correctly manages the group signature

# Proof of Useful Work (PoUW)



- Replace the hash calculation in PoW with “useful” mining work
- Each unit of useful work grants a Bernoulli test
- Similar exponential block time

# REM: Architecture Overview



- Blockchain agent:
  - Collect transactions and generate a block template (a block without PoUW)
  - Publish a mined block to the P2P network and receive reward
- Useful work client:
  - Generate PoUW tasks
- Miners
  - Return useful task results and **(possibly)** provide PoUW

Like Bitcoin, exponential block time is ensured, e.g., 10mins per PoUW among all miners. How?

# Useful Work Metering



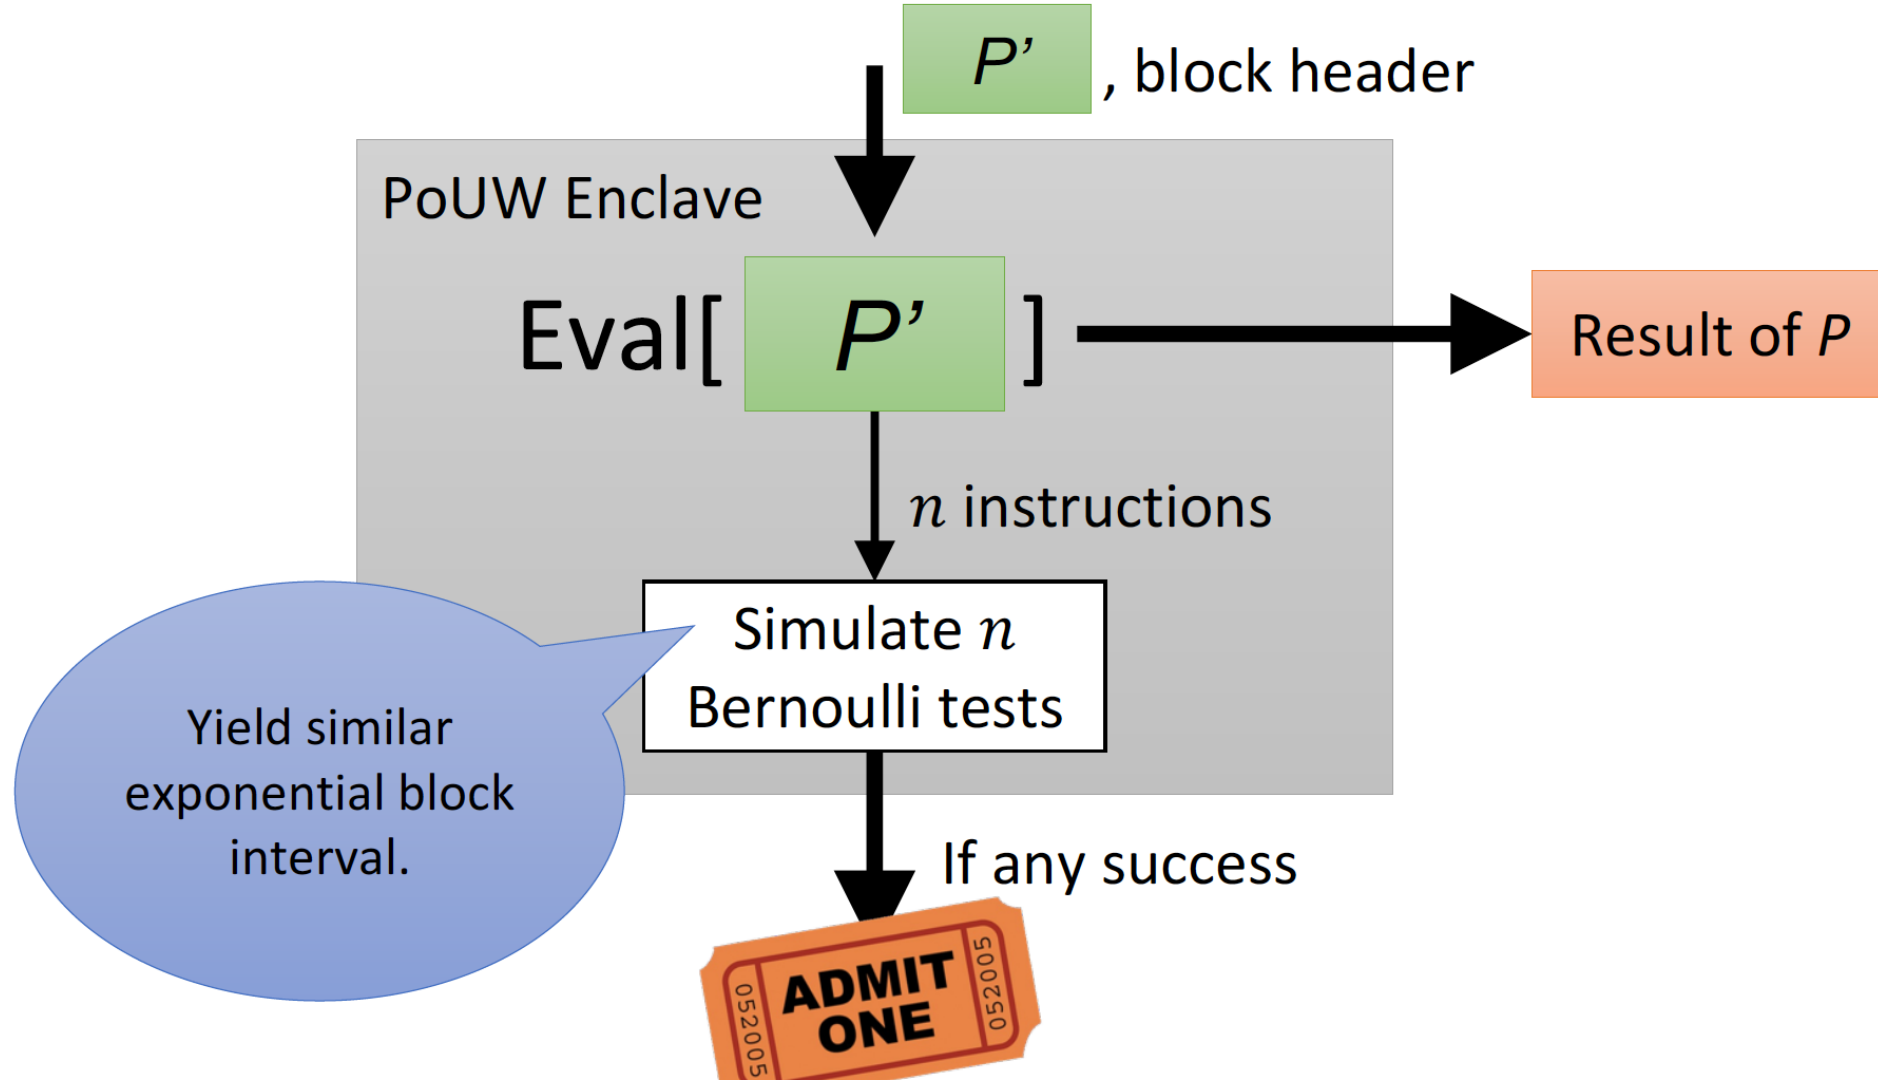
- **Key question 1:** how to meter the effort a miner has conducted?
- Proof of Work: perform two SHA256 hash operations
- REM: perform a CPU instruction that is defined in an PoUW task

# Useful Work Metering (cont.)

- **Key question 2:** how to determine whether an effort is successful, i.e., resulting in a new block?
- Proof of Work: check if the computed hash value is *smaller* than a target value
  - Overall mining effort is measured in terms of the number of executed hashes
- REM: each performed CPU instruction wins a chance to conduct a **Bernoulli trial**
  - Mining times are distributed similar to PoW
  - Overall mining effort is measured in terms of the number of executed useful-work instructions



# REM Miners



# Challenge: Replace PoW with Alternate Resource Lottery

- Other physical resources, with different properties?
  - Disk space
  - Useful computation/storage
  - Etc ...
- What about the coin itself?
  - “Virtual resource mining” --> Proof of Stake (PoS)



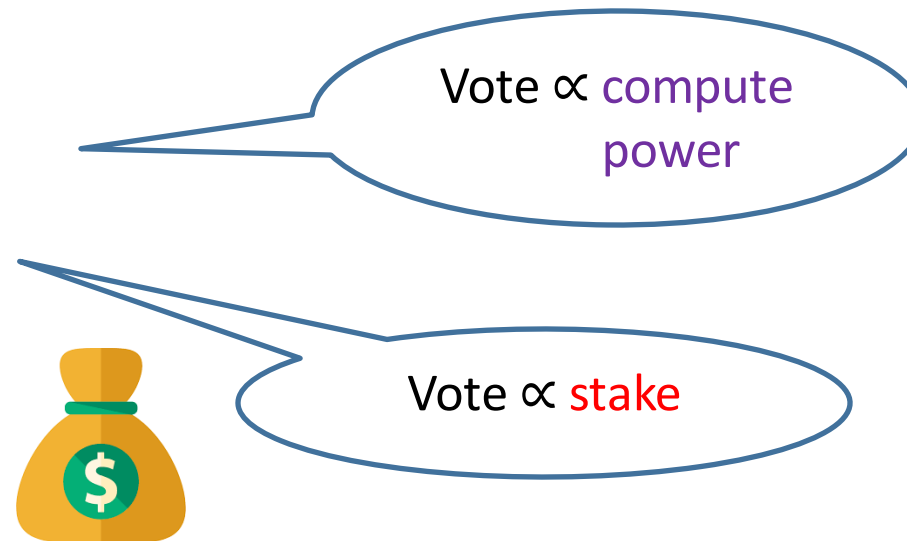
# Challenge: Replace PoW with Alternate Resource Lottery

- Other physical resources, with different properties?
  - Disk space
  - Useful computation/storage
  - Etc ...
- What about the coin itself?
  - “Virtual resource mining” --> Proof of Stake (PoS)



# PoS: Stake-based Lottery

- Blockchain tracks ownership of coins among parties
- Idea: participants elected proportionally to stake  
⇒ No need for physical resources
- Different ideologies
  - Proof-Of-Work (POW)
  - Proof-Of-Stake (POS)



# Ideal PoS Consensus

- Execute in epochs
  - An epoch consists of several time steps
- In each epoch:
  1. Configure a committee of stakeholders
  2. Randomly elect a leader in every time step (based on their stakes)
- In each time step:
  - a. Leader selects the longest available chain
  - b. Leader generates a new block and broadcasts it to the network
- Reconfigure the committee at the end of each epoch
  - To reflect stake changes (due to money transfers)


# PoS: The Challenge

- Difficulty: Electing a coin (owned by a stakeholder) requires randomness
- If the adversary can bias the randomness, things get difficult



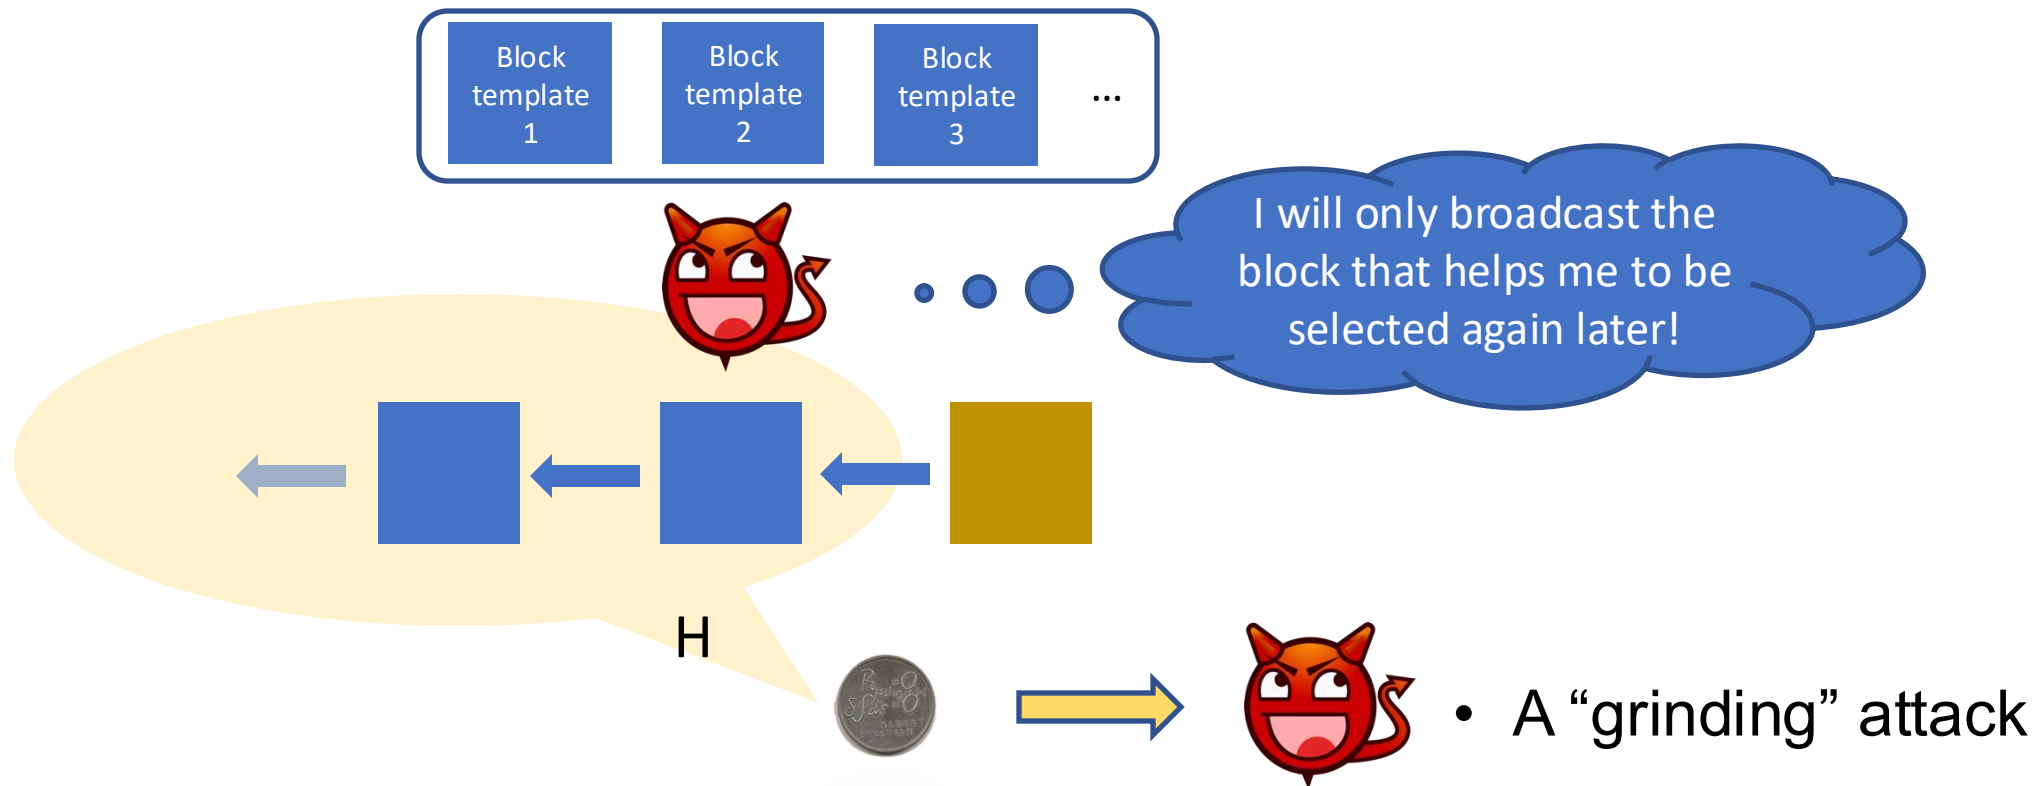
# An Idea

- Blockchain already contains effectively random data (e.g., block hashes)
- Obvious approach:  
Just hash the current blockchain!



# An Idea ... confounded!

- However, adversary can try to bias the randomness in their own favor





# PoS Designs “in the wild”

- PoS blockchains have appeared “in the wild”:
  - NXT
  - Peercoin
  - DPoS (BitShares, Steem, EOS)
  - Casper (Ethereum)
  - Etc...

# PoS Designs with Rigorous Guarantees

- Eventual (Nakamoto-style) Consensus:
  - Ouroboros
  - Ouroboros Praos
  - Ouroboros Genesis
  - Snow White
- Block-wise Byzantine Agreement:
  - Algorand

Here we will only elaborate two PoS designs with rigorous security guarantees:

Ouroboros: A Provably Secure Proof-of-Stake Blockchain Protocol, Aggelos Kiayias et al., *in Proc. of CRYPTO*, 2017

&

Algorand: Scaling Byzantine Agreements for Cryptocurrencies, Yossi Gilad et al., *in Proc. of SOSP*, 2017

# Ouroboros: 10000ft View

- Assumes **synchronous** time & communication
- Guarantees
  - persistence: stable transactions immutable
  - liveness: new transactions included eventually


IF

- Adversary has minority stake throughout
- Adversary subject to corruption delay
  - Corrupt an honest stakeholder after some delays




# Ouroboros: Overview (1)

- Assuming root-of-trust
  - Initial stake distribution of stakeholders is hardcoded in the genesis block
  - Trusted to achieve honest majority



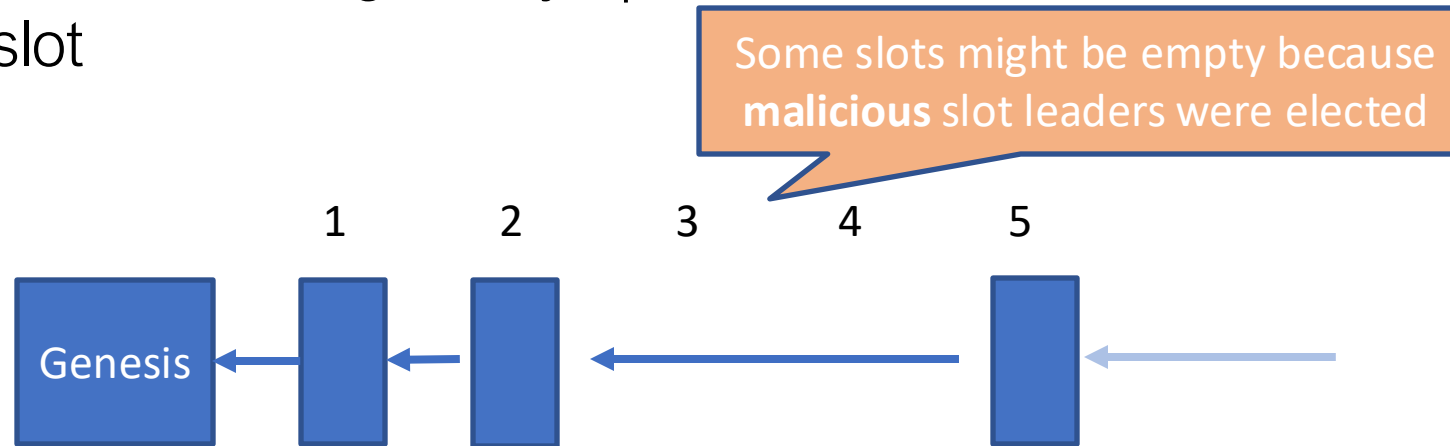
# Ouroboros: Overview (2)

- Randomly elect leaders in proportional to their stake
  - Leveraging coin-tossing and verifiable secret-sharing to generate randomness for the next epoch (will introduce later)
  - The generated randomness and stake distribution define  $R$  slot leaders in the next epoch
    - Randomly select coins in the previous epoch, and elect the corresponding **coin owners as slot leaders** in the next epoch
      - More coins (money) you have, more chances to be elected!



# Ouroboros: Overview (3)

- Blockchain validation:
  - a. Starts with genesis block
  - b. A sequence of blocks follow, associated with increasing slot numbers
  - c. No conflicting transactions
  - d. Each block signed by  $L_i$ , the leader associated with that slot




# Randomness Generation

- Recall: randomness is needed for next epoch's leader election
- Idea:
  - Use a secure coin-tossing protocol to generate it during the previous epoch
- Result:
  - If a majority of the leaders in the previous epoch are honest, we can get clean (unbiased) randomness for next epoch



# Coin-tossing Protocol

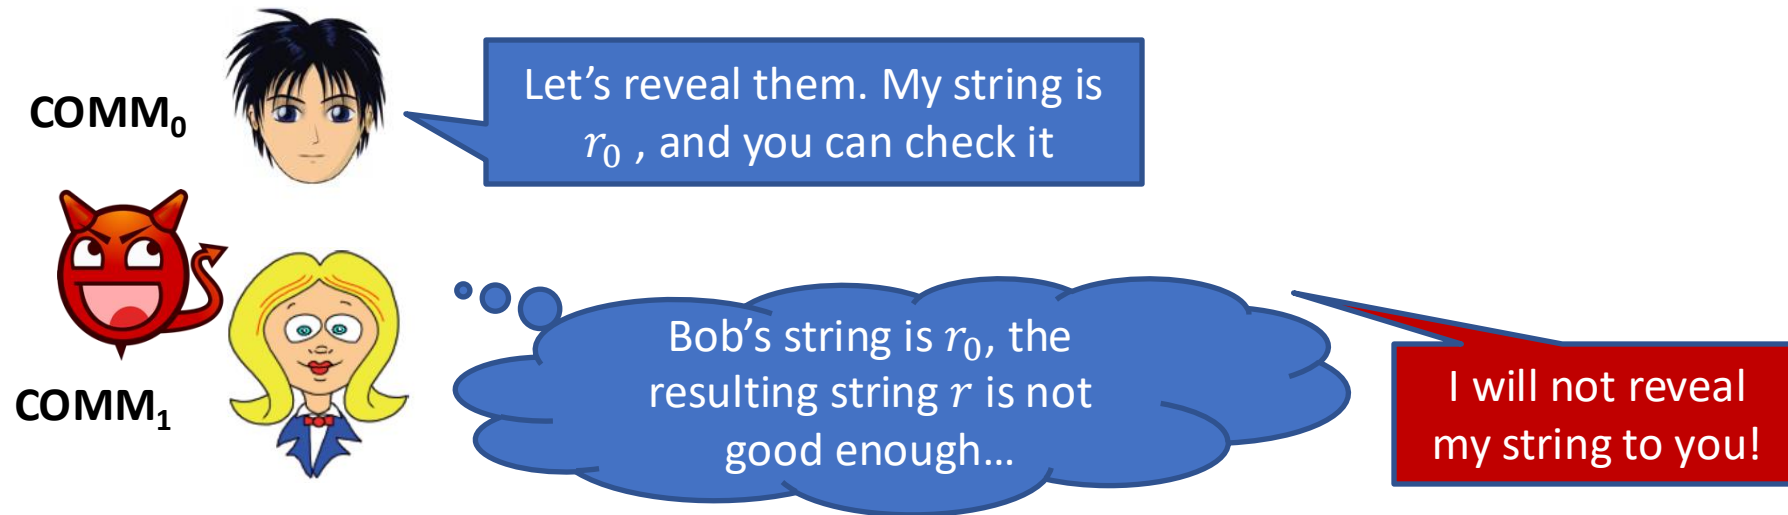
- Goal: generate an **unbiased** random string in a distributed setting
- We start with strawman design ①:
  - Simply outputs  $r = \bigoplus_{i=0}^{n-1} r_i$
  - Here, we demonstrate the case that  $n = 2$



What is the value of  $r$  now?

# Second Attempt

- We can have each peer **commit to their chosen input** before seeing other inputs
  - Commit-then-reveal approach
  - Assume Bob has committed its string  $r_0$  as  $\text{COMM}_0$ , and Alice has committed its string  $r_1$  as  $\text{COMM}_1$



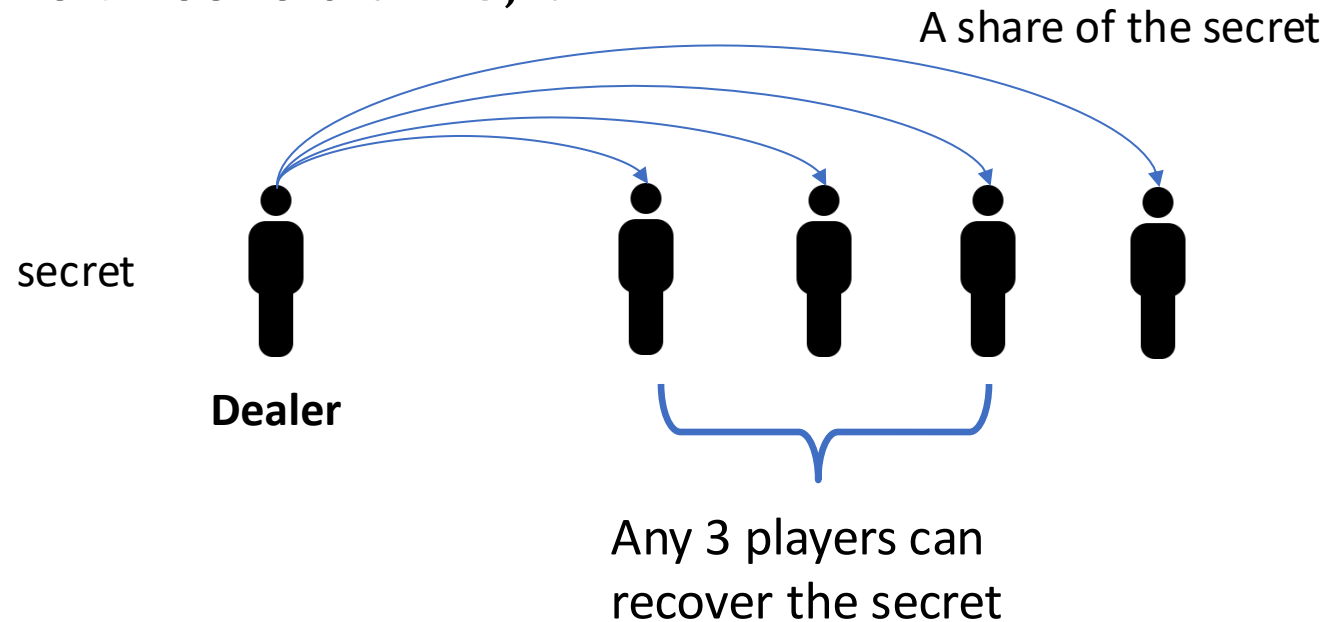
Second attempt, failed!

# Third Attempt

- We wish to ensure that a **dishonest peer** cannot force the protocol to **abort** by refusing to participate
  - We need more than 2 players!
  - Leverage a  $(t, n)$ -secret sharing scheme
    - $n$  is the number of players
- $(t, n)$ -secret sharing
  - A secret is shared among  $n$  persons, and we only need  $t$  shares to recover it
  - E.g., Shamir's secret-sharing scheme
    - Can extend to achieve public verifiable to address bad shares that prevent correct recovery of the shared secret


# Simplified Demonstration

- Consider a  $(3,4)$ -secret sharing scheme for 5 players, with one player acted as the dealer (for sharing a secret)
  - The threshold  $t = 3, n = 4$



# Simplified Demonstration (Cont.)

- Each player will act as the dealer once and share a secret to others
  - Thus in the case of 5 players, each player will eventually have 4 shares of secrets



# Simplified Demonstration (Cont.)

- Reconstructing all secrets (i.e., 5 secrets in our previous example) and XORing all secrets to generate the final random string
- Analysis:
  - As long as the majority of players (e.g., 3) are honest, all secrets can be correctly recovered
    - Since we set  $t = 3$  as our threshold in the secret sharing scheme
  - A random and unbiased string is guaranteed!

# Protocols for Augmenting Ouroboros

- **Ouroboros Praos** with improved assumptions:
  - In a **semi-synchronous communication** model
  - Despite fully adaptive corruptions
    - I.e., attacker now can corrupt honest player that has been elected in previous epochs
- **Ouroboros Genesis:**
  - Global UC formalization
  - Improved bootstrapping procedure
    - New chain selection rule (instead of longest chain only)
  - Dynamic availability
    - Achieves security with node join and leave


Algorand: Scaling Byzantine Agreements  
for Cryptocurrencies, Yossi Gilad et al.,  
*in Proc. of SOSR, 2017*

**(not to be included in final exam)**



# Algorand: The View from Mars

- Each new block in Ouroboros is generated by **one stakeholder**
- What if:
  - We replace that specific stakeholder with **a committee**? => Algorand



# Algorand: Overview

- Ingredients:
  - Committee formation: Cryptographic Sortition
    - Verifiable random functions (VRF)
  - Consensus: new Byzantine Agreement protocol (**BA\***)
- Assumption:
  - Adversary can only control **less than 1/3** of total money

Question 1: how can we securely elect a committee from the entire network?

=> Algorand uses a technique called  
cryptographic sortition


# Cryptographic Sortition

- Goal: Select a subset of users
- **Simple idea:** Get the users to interact and talk amongst themselves to elect a committee
  - Is it secure?
- However, the adversary could be listening to the election process
  - Might **target some users** before they can participate in the interactive protocol (denial of service attacks!)



# Cryptographic Sortition in Algorand


- Goal: Local and non-interactive way to select a subset of users (*forming a committee*)
  - Implemented using Verifiable Random Functions (VRF)



Hash is **pseudo-random** and **distributed uniformly between  $[0, 2^{\text{hashlen}-1}]$**

# Result: A Verifiable Hash Value

- Given  $\langle \text{Hash}, \text{PK}, \text{proof} \rangle$ , anyone can verify the validity of Hash
- Since Hash is uniformly distributed between  $[0, 2^{\text{hashlen}} - 1]$ 
  - $T = \text{Hash} / 2^{\text{hashlen}} \in [0, 1)$



**Since Hash is verifiable, T is verifiable too!**

Next important property:



=> Making the chance that each stakeholder is selected **proportional to** its stake in the blockchain!

# Background: Binominal Distribution

- $B(k; w, p)$ : Probability of getting  $k$  successes in  $w$  trials, where probability of success in each trial is  $p$

$$B(k; w, p) = \Pr(X = k) = \binom{w}{k} p^k (1 - p)^{w-k}$$

for  $k = 0, 1, 2, \dots, w$ , where

$$\binom{w}{k} = \frac{w!}{k!(w-k)!}$$

is the **binomial coefficient**

- One useful property:

$$\sum_{k=0}^w B(k; w, p) = 1$$



# Notations in Algorand

- $W$ : current total amount of money units in the system
- $\tau$ : threshold, denoting expected number of money units selected
- $p$ :  $\tau/W$
- $w$ : stake/money of a user

# $B(k; w, p)$ Meaning in Algorand


- If  $p$  represents the probability that a coin is selected
  - $\Rightarrow B(k; w, p)$  defines the probability that exactly  $k$  coins are selected from the stakeholder's  $w$  coins
  - $\Rightarrow$  For example, Bob has **2 coins** and the probability  $p$  is pre-defined
    - $B(0; 2, p)$  is the probability that **no coin** from Bob is selected
    - $B(1; 2, p)$  is the probability that **one coin** from Bob is selected
    - $B(2; 2, p)$  is the probability that **two coins** from Bob are selected



$$\sum_{k=0}^2 B(k; 2, p) = 1$$

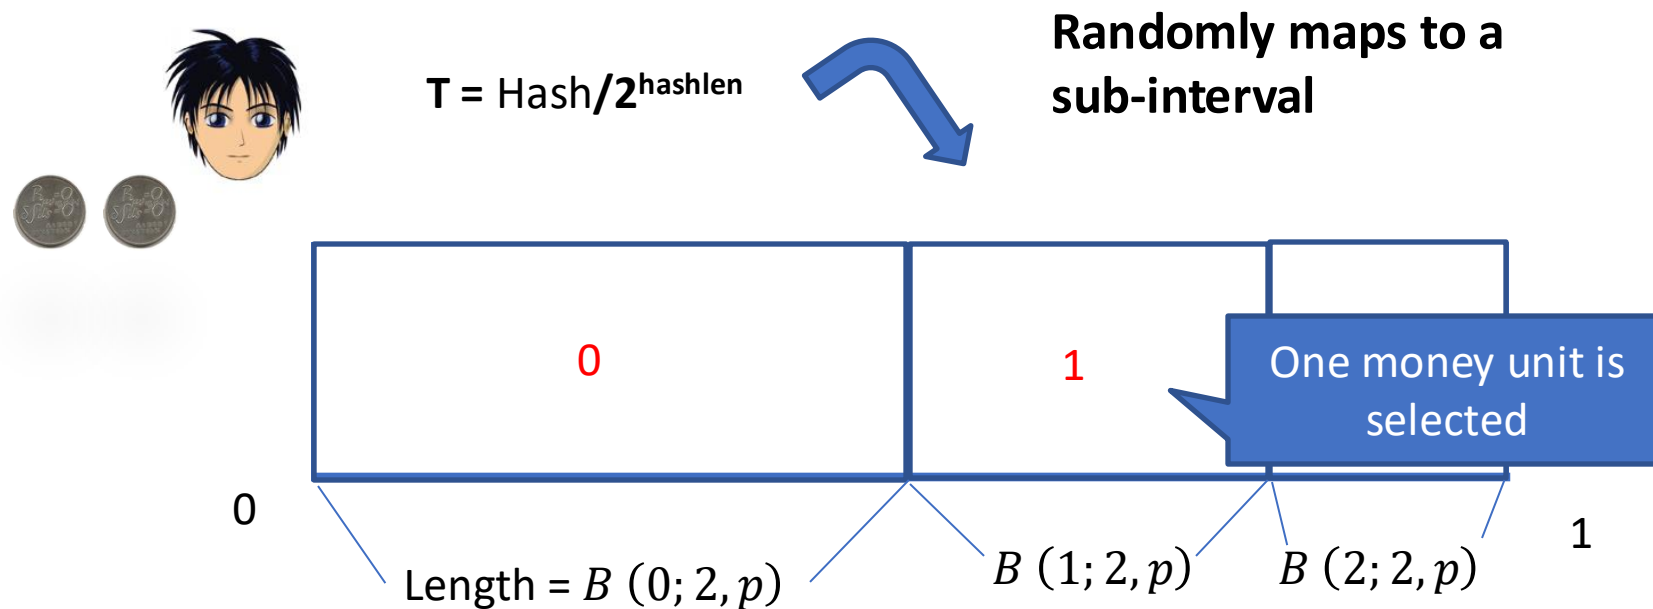
# Achieving Stake-based Selection (1)

- Idea: for each stakeholder, we divide the interval  $[0,1)$  into  $w + 1$  sub-intervals, where  $w$  is the # of money units the stakeholder owns
  - Requirement: each sub-interval  $j$  still reflects the probability of  $B(j; w, p)$
- With 2 coins (money units)



# Achieving Stake-based Selection (2)

- We have a verifiable value  $T = \text{Hash}/2^{\text{hashlen}}$  that uniformly distributed in interval  $[0,1)$



# Achieving Stake-based Selection (3)

- Observation: more money units might be selected at a **richer** stakeholder
- $\tau$  can be adjusted to control the # of (totally) selected money units
  - Since  $p := \tau/W$
- Definition: IF more than one money unit is selected, the corresponding stakeholder is selected in the committee

# Result: A Committee with Different Voting Weights

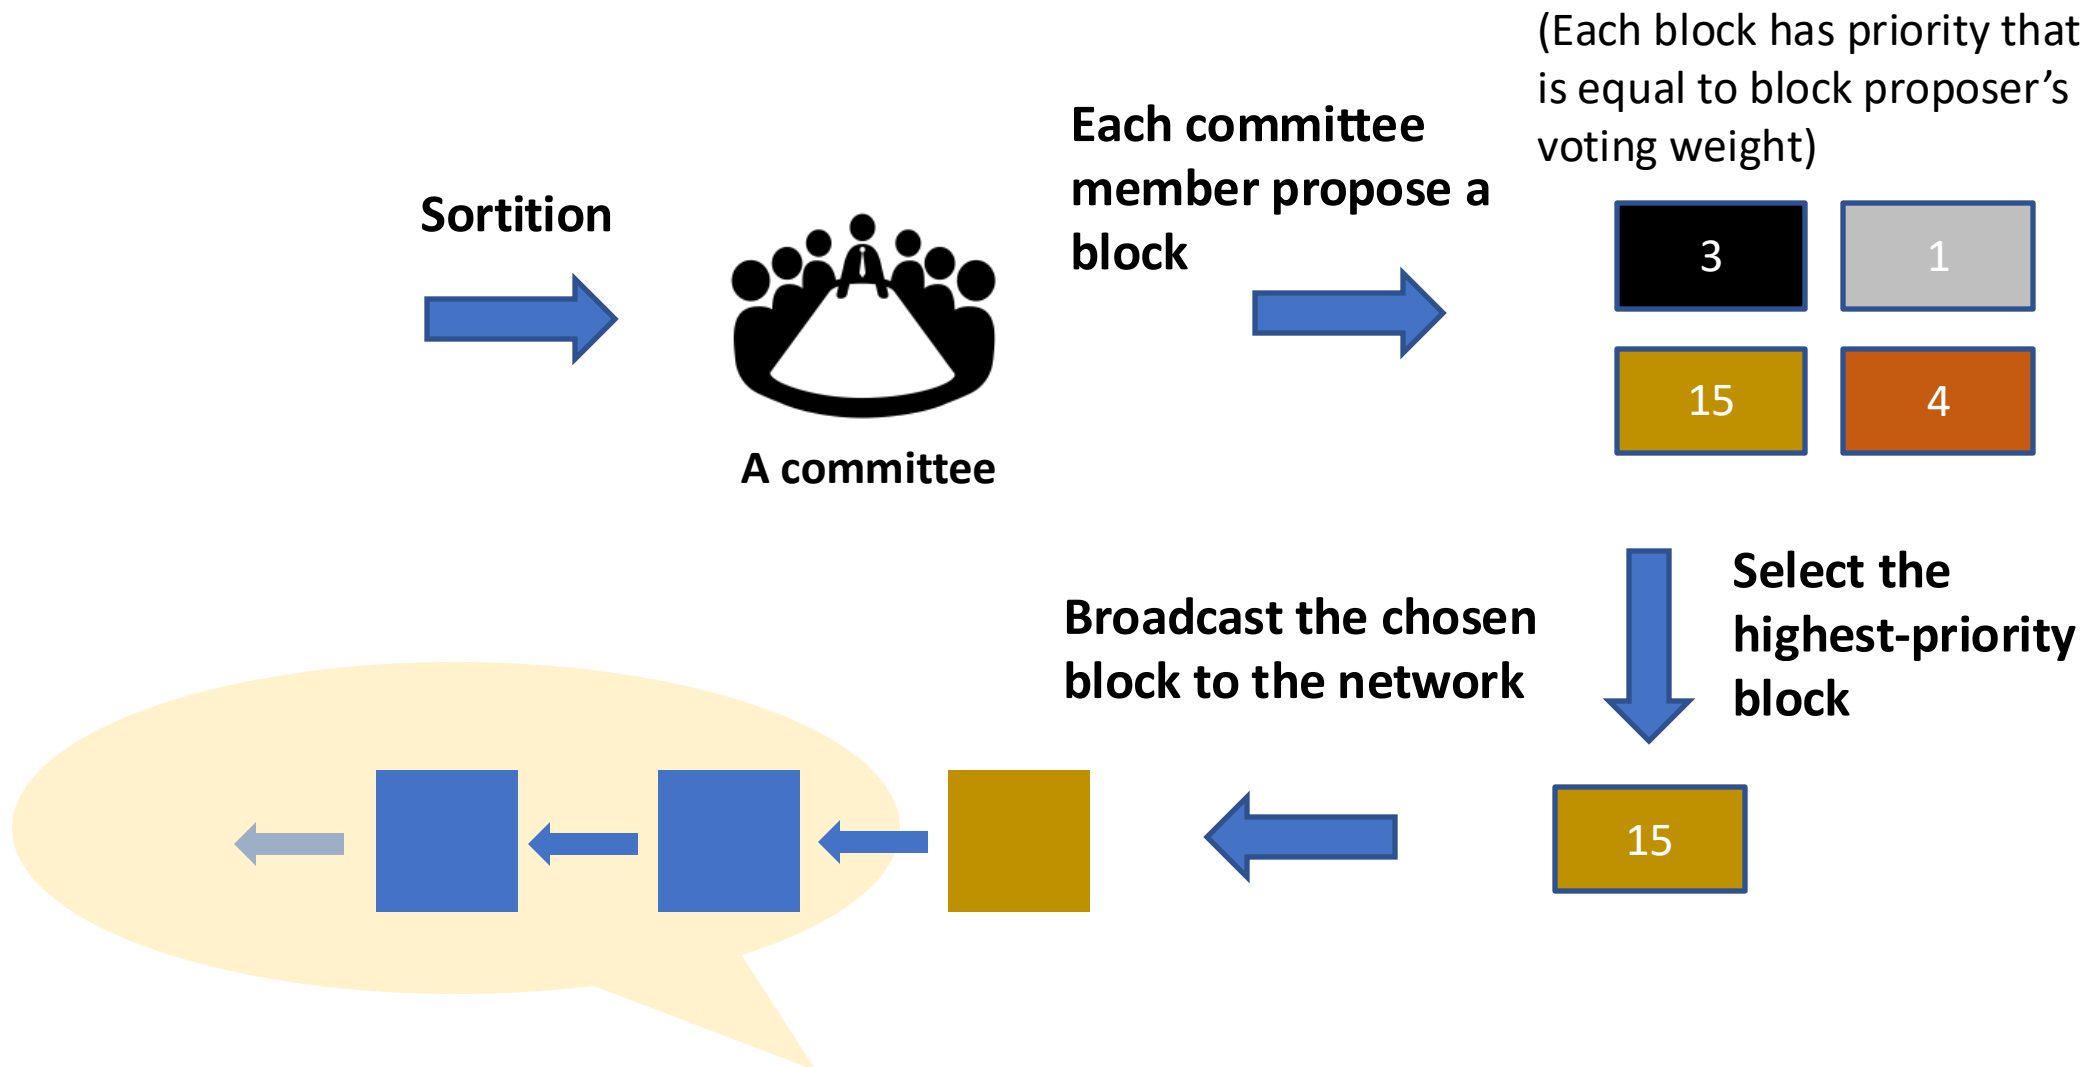


Identity	# of selected money units
Alice	4
Bob	1
Carol	15
John	3
Peter	9

Vote weight

Next question: How to generate and commit new blocks with this committee?

# Strawman Design



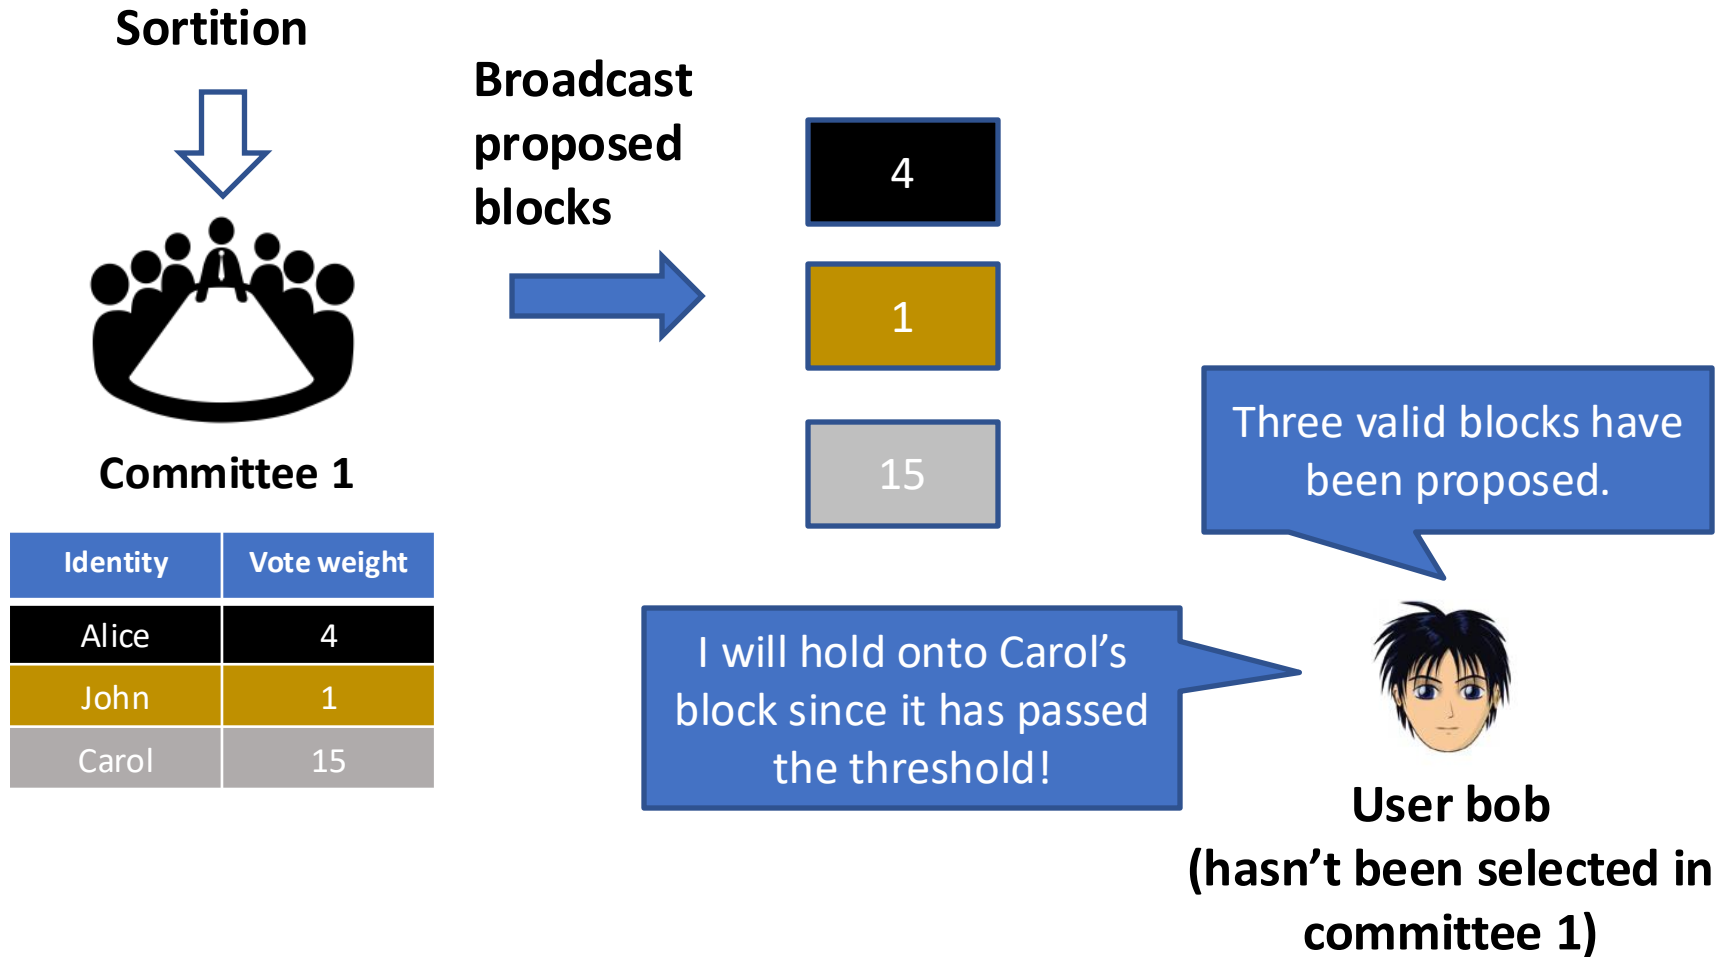
# Actually, Algorand Consensus is More Complicated

- In each step, a **different** committee is elected through cryptographic sortition
  - Address attacker who can corrupt some honest committee members after the election
- Basic strategy: weighted majority voting:
  - Each committee member will **broadcast** their **vote** for their block
    - Vote for **highest priority** block
    - All users can see this message
  - Users that receive more than a **threshold** of votes for a block will **hold onto** that block



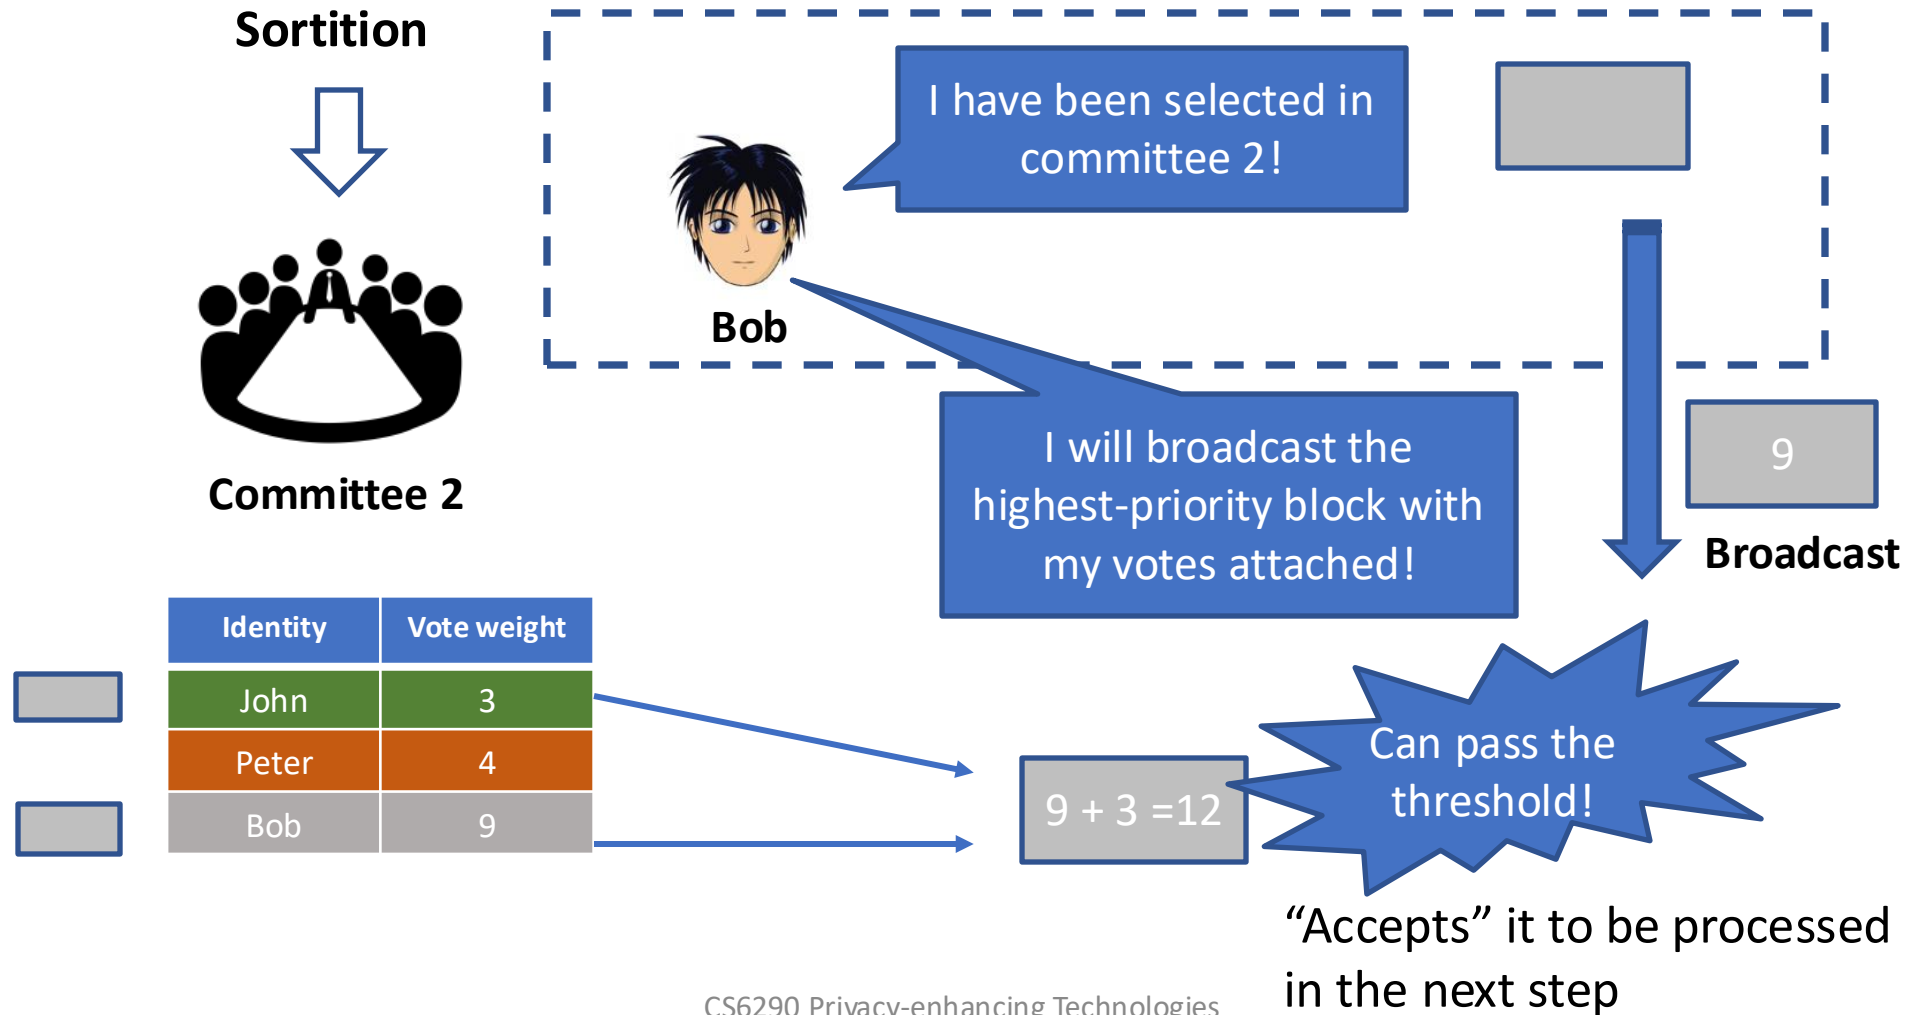
# Simplified Demonstration (1)

(threshold=10)



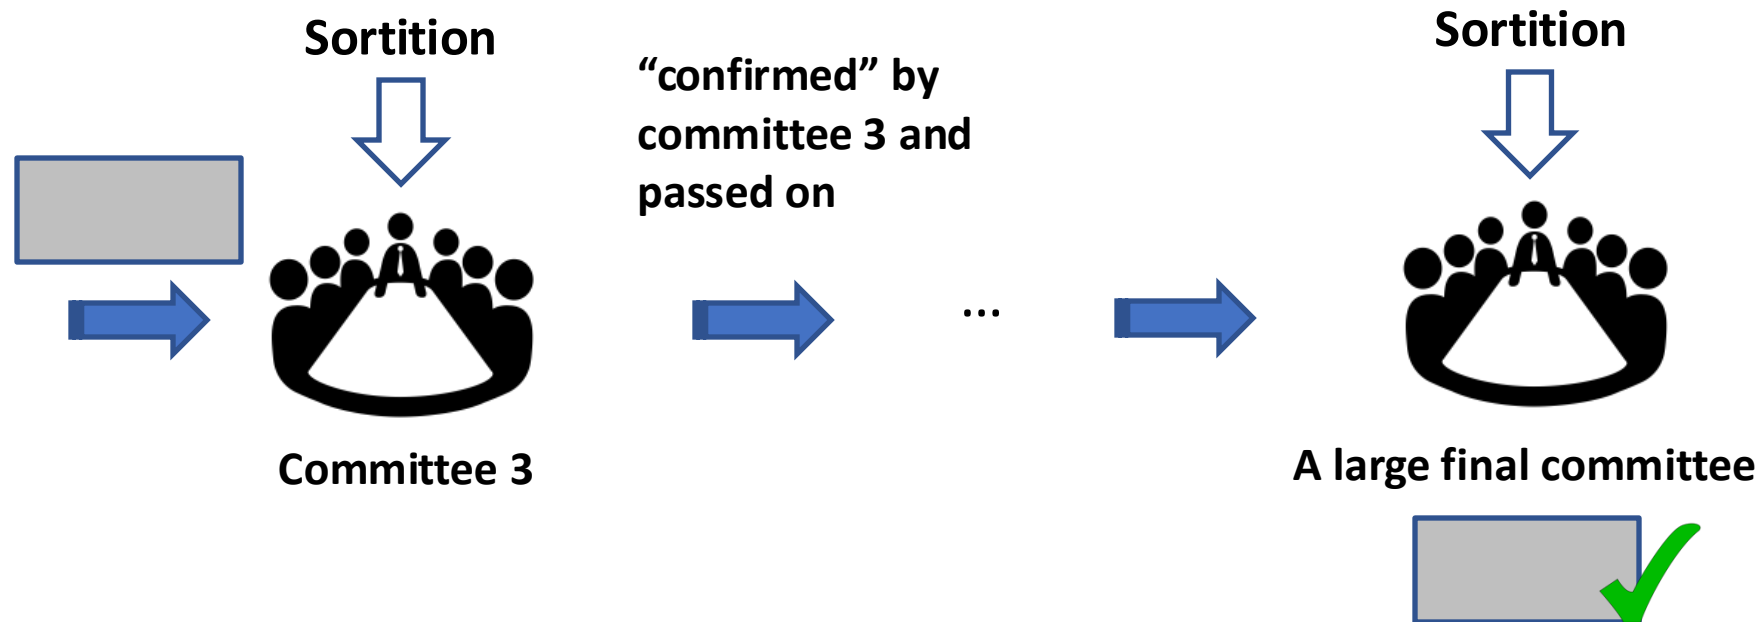
# Simplified Demonstration (2)

(threshold=10)



# Simplified Demonstration (3)

- So on and so forth ...
- Finally, being confirmed by a final committee and committed to the network




# Performance Highlights

- BA\* protocol in expectation terminates
  - In 4 steps
    - in the case where the block proposer with highest vote is “honest” and the network is strongly synchronous
  - Or in 13 steps
    - in “disaster” case
    - Might require additional algorithm for recovery
- ~1 min to confirm transactions vs an hour (6 blocks) in Bitcoin


# Wrapping Up

- Ouroboros uses **one stakeholder** for handling block proposal
- Algorand (in each step) maintains **one committee** for handling block proposal



Can we go further?  
Say, with multiple committees working  
in parallel?


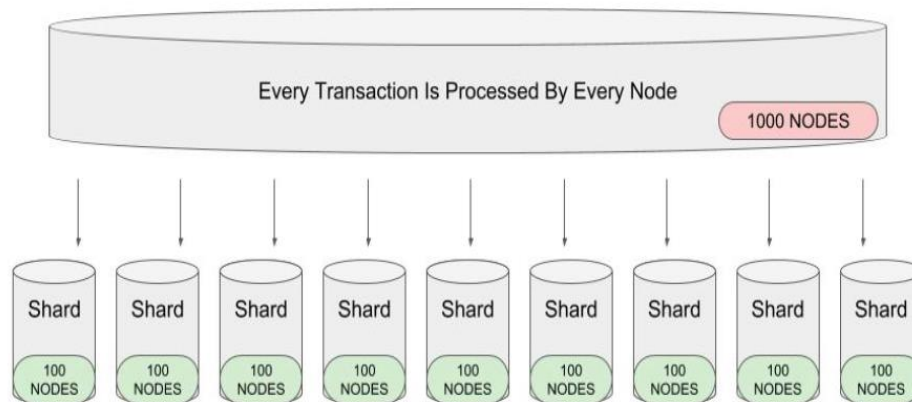
# Sharding: A Glimpse



A blockchain system that scales with the number of players participated in the system

# Blockchain Sharding Protocols

- Motivation: single committee improves performances, but it does not **scale out**
- Idea:
  - **splitting transactions** among multiple committees (shards) which process these transactions in parallel






# But Scaling out Blockchains is Not Easy



Figure credit to Omniledger

# The Scalability Trilemma (Vitalik)




E. Kokoris Kogias et al., Enhancing Bitcoin Security and Performance with Strong Consistency via Collective Signing, USENIX Security 2016

# An Overview of Existing Blockchain Sharding Methods

- Main components:
  - Miner assignment + intra-shard consensus + security refinements
- UTXO-based mode:
  - Elastico CCS'16, Omniledger S&P'18, Rapidchain CCS'18, Monoxide NSDI'19
  - Projects: Zilliqa, Blockclique, ...
- Account-based mode:
  - Chainspace NDSS'18, DANG et al. SIGMOD' 19


# Miner Assignment

- “Buy ticket” before participating
  - Ticket here can be a Proof-of-work (Rapidchain, Omniledger) or a stake-based VRF proof
  - Mainly to address *Sybil attacks*



# Shard Configurations


- The number of shards is defined in advance
- Each shard has a shard\_ID (can be a number)





# Which shard a miner should go?

- Let miners choose?
  - All malicious miners can choose the same chain
- Randomly assign miners?
  - Preserve security for adequately large shard size

X validators building one chain.  
Need to corrupt 0.51x




X validators building 10 chains  
Need to corrupt 0.051x



# Transaction Assignment

- Create **disjoint transaction sets**
  - Hash the ID of the incoming transaction
  - See which shard it maps to and forward it to that corresponding shard



# Main Confronted Challenges

- **Intra-shard consensus**
  - How to reach consensus in each shard
- **Cross-chain verification**
  - Inputs/outputs of a transaction might be maintained in different sharded blockchains



# Intra-shard Consensus

- Random miner assignment guarantees each shard to have an upper of malicious miners
  - 1/3 in Omniledger, Chainspace and Elastico
  - 1/2 in Rapidchain
- As the miners in each shard are identified, all existing standard consensus protocols can be used
  - PBFT is used in Omniledger, Chainspace and Elastico
  - An enhanced consensus algorithm (with  $\frac{1}{2}$  resilient) is used in Rapidchain

# Final Consensus?

- In *Elastico*, at each epoch, the blocks generated by shards will be checked by a final committee
  - *For making each shard small (100 miners)*
  - Only if the final committee confirms the blocks, they become final at the global state (slower)



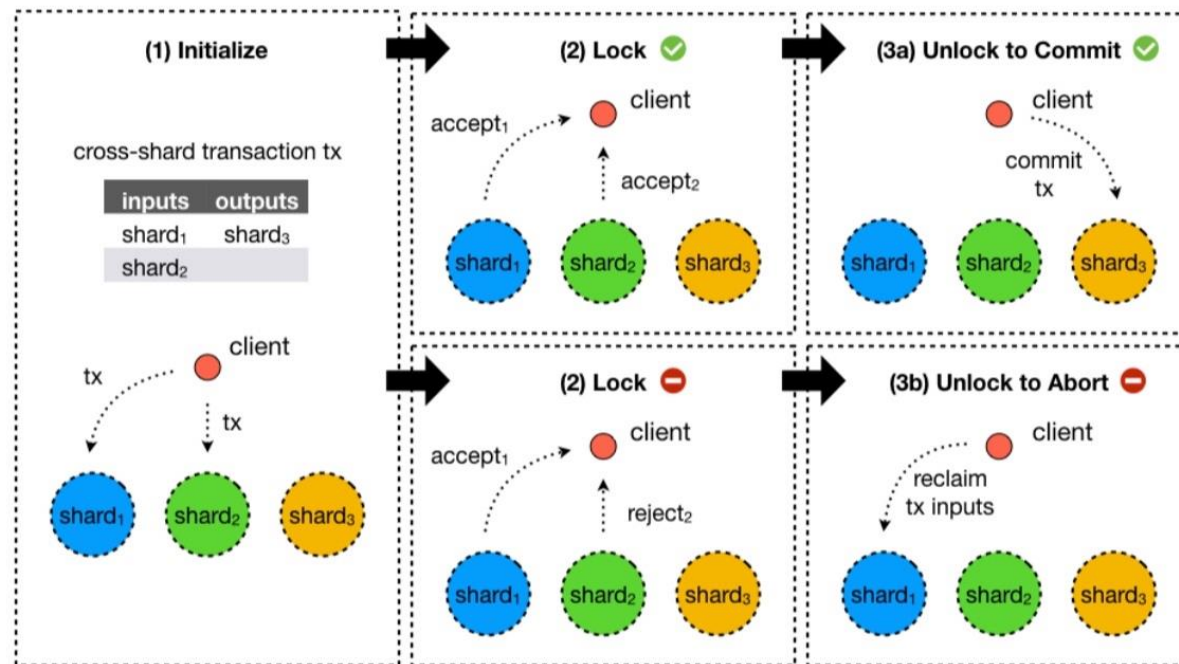
*Subsequent works (e.g., Omniledger and Rapidchain)  
remove the final committee by increasing the size of shards*

# Cross-shard transaction

- The following example is courtesy of Andrew Miller:
  - A user wants to purchase a train ticket and reserve a hotel
  - The operation should be atomic: either both reservation succeed or neither do
- A (in shard\_1) , B (in shard\_2) -> C (in shard\_3)
- Both A, B should be accepted before C happens

# Cross-shard transaction

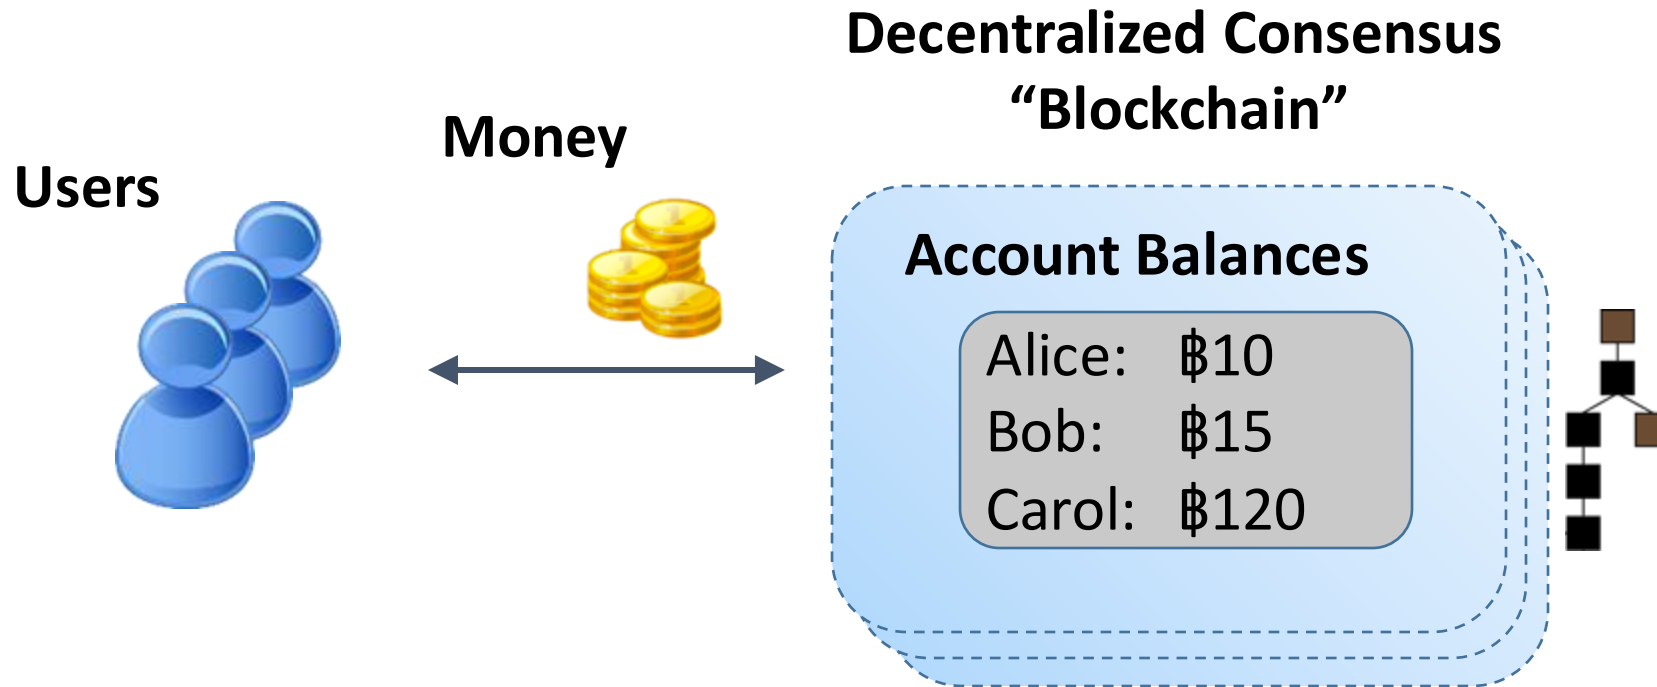
- If the two objects are situated in two different shards, cross-shard operation should be involved
  - **Atomic commit:** “lock-and-free”, assuming synchronous communication



Challenge: what if the communication between two shards is *asynchronous*; how to *reduce the overhead*

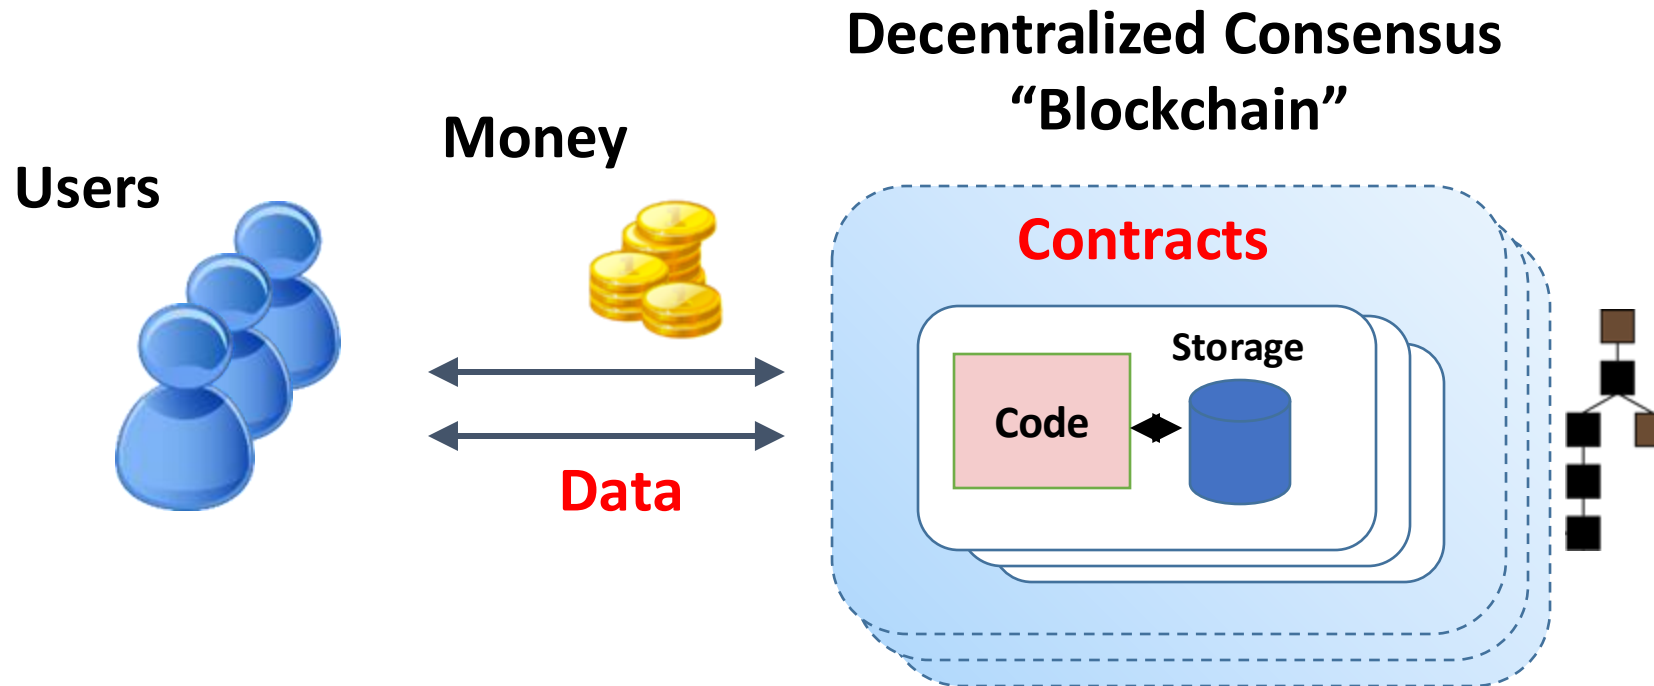
# Recall: Bitcoin

- Just one application on top of a blockchain: token exchange



# Smart Contracts

- User-defined programs running on top of a blockchain



# What's a (decentralized) smart contract?

- Executable object on blockchain
- Scripted in Turing-complete language
  - Bitcoin has highly restricted language
- Code defines contract, e.g.,
  - Financial instrument
    - If GOOG rises to \$1,000 by 30 June 2015, assign 10 shares from Alice to Bob and pay Alice \$10,000
  - Szabo (1997): Smart contract reassigns physical access to your car from you to your bank if you don't make a payment
- Autonomous: *Enforced by network*



Can we build a blockchain that  
natively support smart contracts?

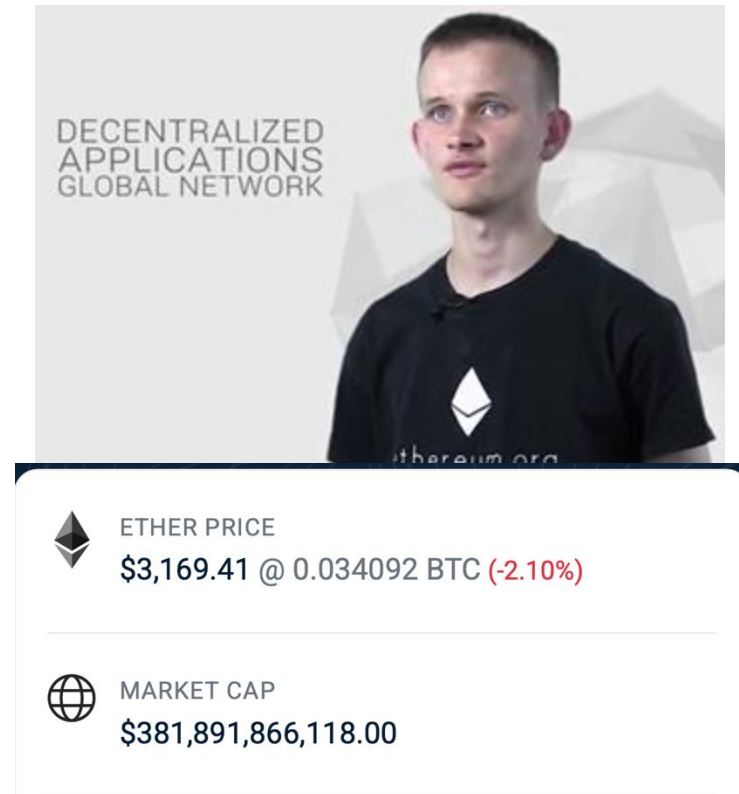
=>Ethereum






# About Ethereum

- Founded by *Vitalik Buterin*, Gavin Wood and Jeffrey Wilcke in 2014
- Core: Ethereum Virtual Machine (“EVM”)
- “general purpose computation” programming language
- Every node of the network runs the EVM and executes the **same instructions on blockchain**.
- Sometimes described evocatively as a “world computer”



# Introducing the Ethereum platform




# In-built cryptocurrency



- Ether
  - Mined in a way similar to Bitcoin
  - Block reward: ~ 3 ETH
- Currently \$3100+ per ether
- The crypto-fuel for the Ethereum network
  - Healthy ecosystem: a form of **payment** to compensate the machines for executing the requested operations
  - Also, an incentive ensuring that developers write quality applications (wasteful code costs more)

# General purpose script language

- High-level languages :  
Serpent, *Solidity*, LLL
- Similar to Java script
- Written contracts are compiled to *EVM code* and deployed on EVM
- Each contract has unique address on EVM
- Execution?

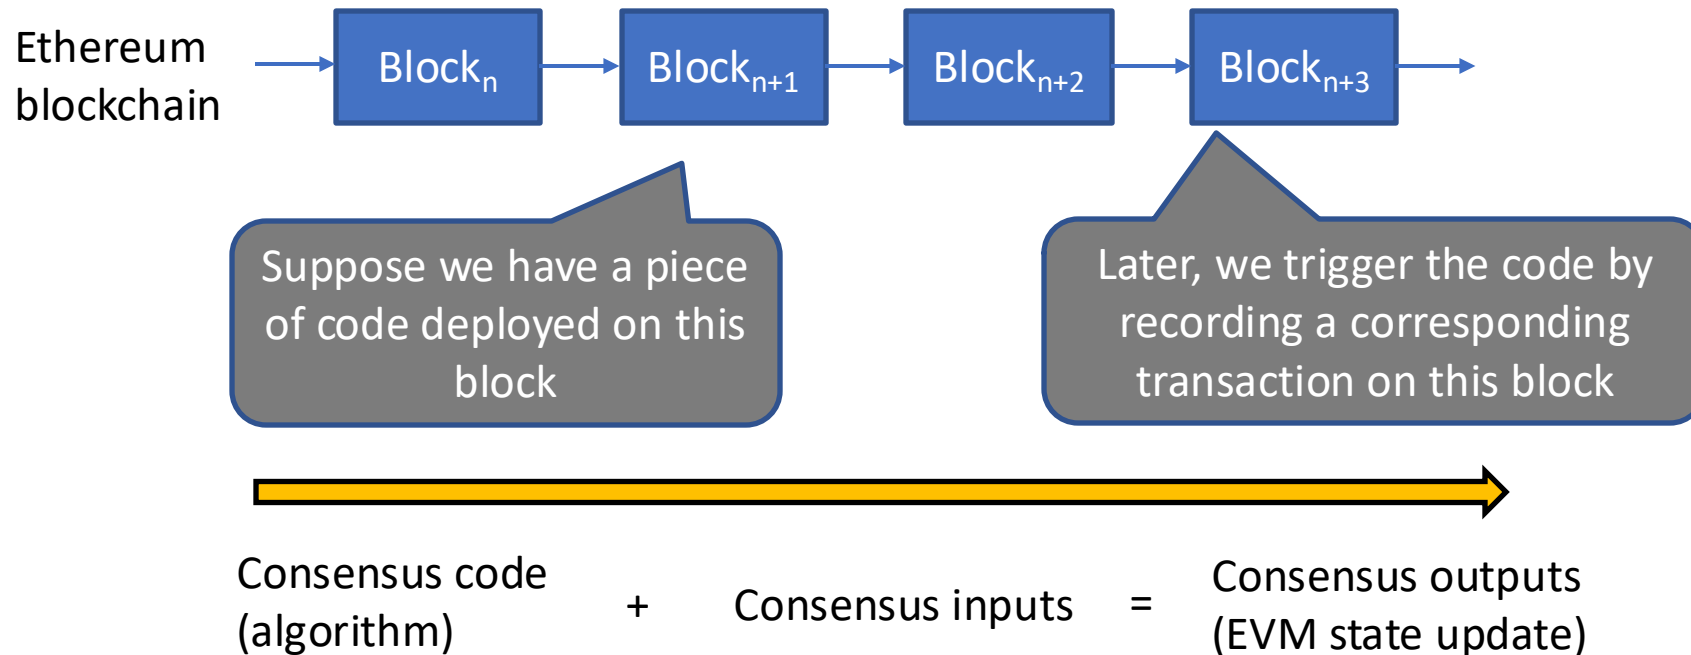


# Code execution

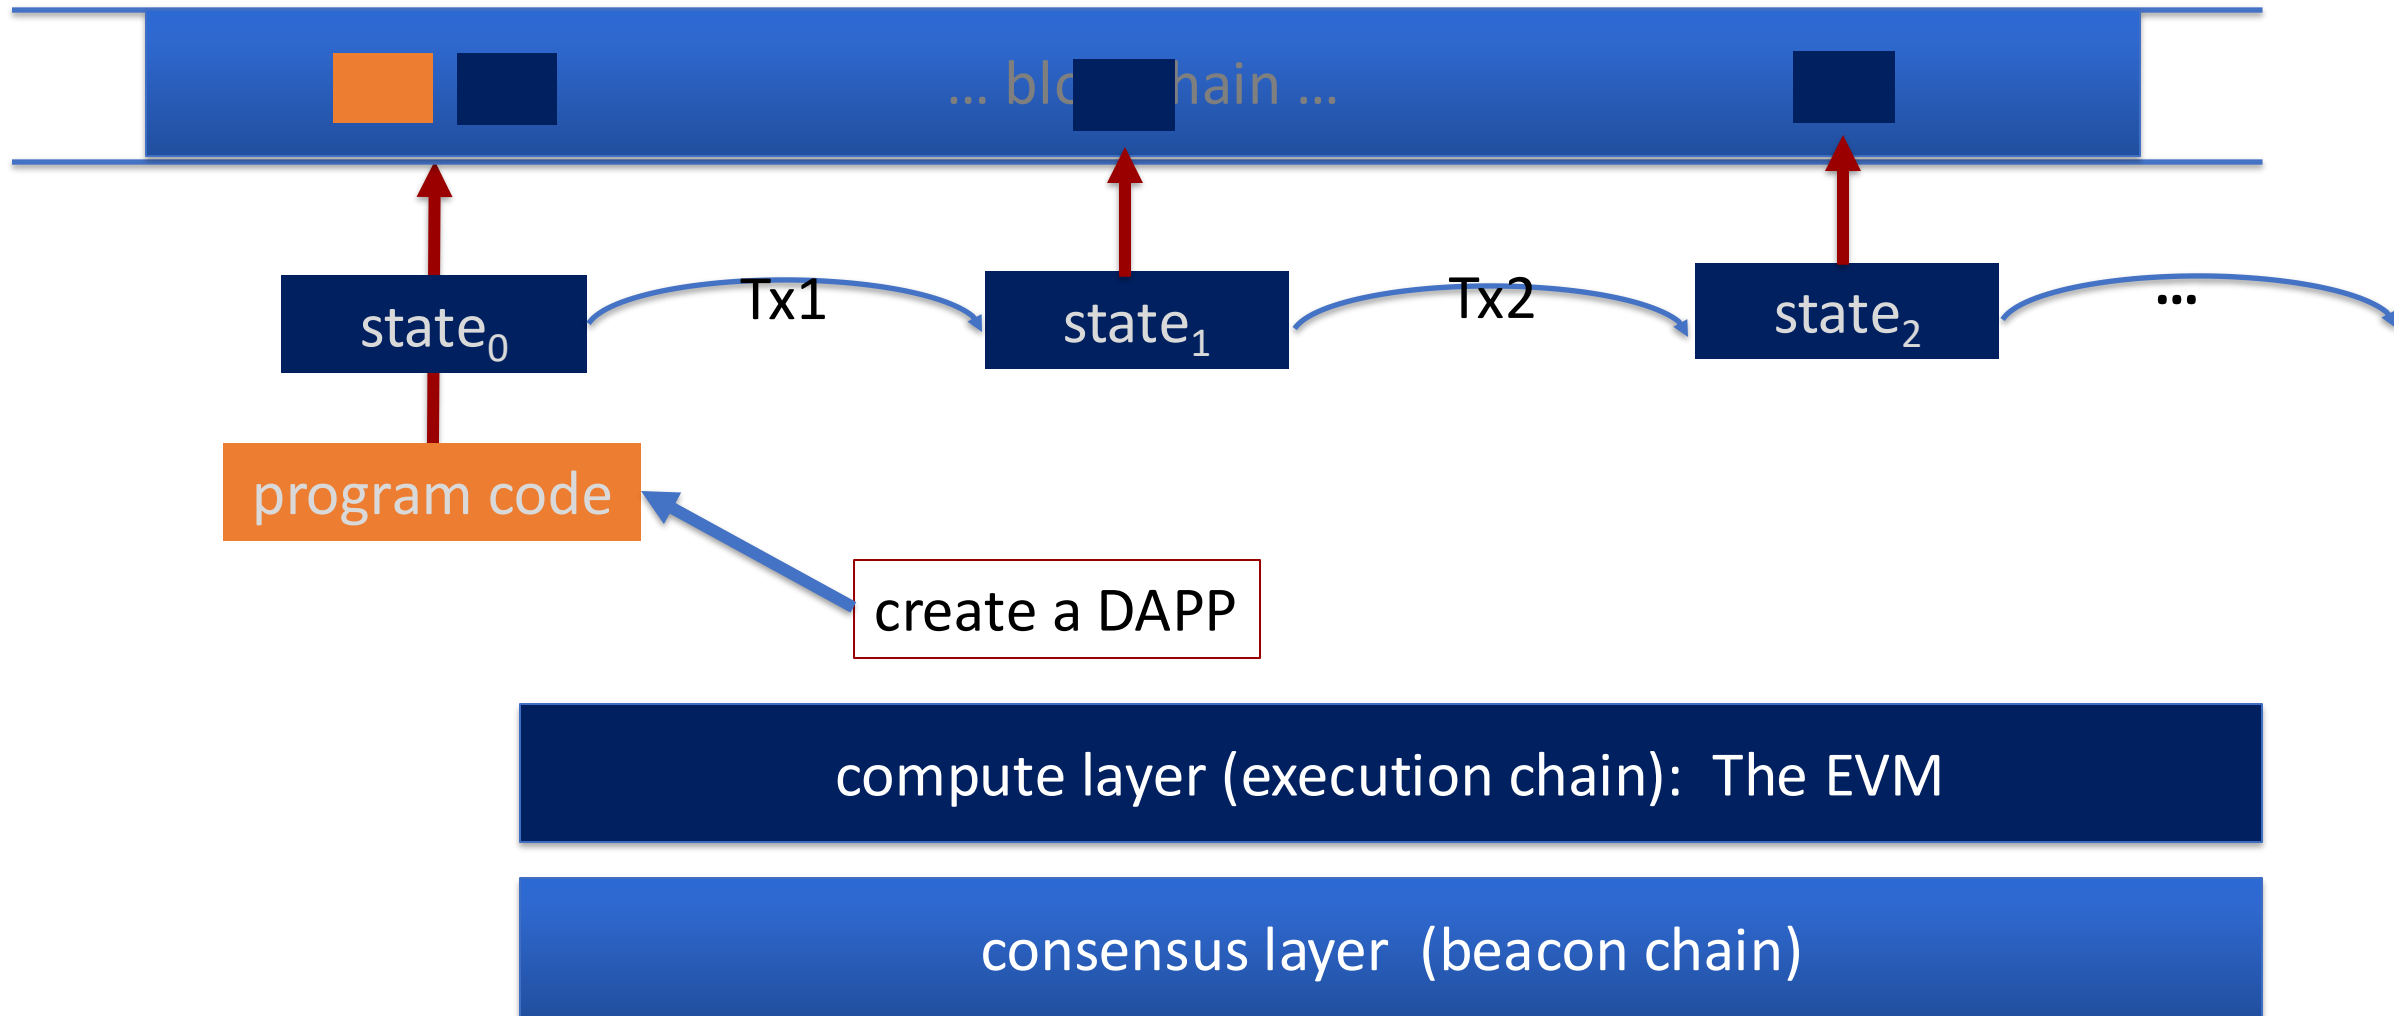
- Deployed contracts (code) are triggered by transactions
  - Transactions specify a TO contract address it sends to
  - Record on the blockchain
  - EVM executes it with the deployed code
- Code can:
  - Send ETH to other contracts / individuals
  - Read / write storage
  - Call (i.e., start execution in) other contracts

# Code execution

- Executed by all nodes in the Ethereum network
  - Ask any node to learn the updated state of the EVM



# Running a program on a blockchain (DAPP: decentralized applications)



# Gas system

- We cannot tell whether or not a deployed code will run **infinitely**
  - Waste computing resources
  - Might damage system robustness
- Ethereum solution: charge fee per **computational step (gas)**
  - e.g. adding two numbers costs 3 gas, calculating a hash costs 30 gas, and sending a transaction costs 21000 gas
  - Special gas fees also apply to *storage* operations
  - Currently, the gas price is around  $2 * 10^{-9}$  ether (<https://ethgasstation.info>)



# Gas limit

- Specify the *maximum amount of gas* the sender is willing to buy in a transaction
  - The code processing stops If the gas used exceeds this limit during execution
  - The sender still has to pay for the performed computation
- Gas limit also applies to each block
  - Similar to the maximum block size in Bitcoin
  - Voting mechanism (can upvote/downvote gas limit by 0.0976% every block)

# Writing Ethereum Smart Contract

- Programming language: Solidity
  - Contract-oriented
  - Syntax similar to Javascript
- A contract is similar to a class
  - State variables
  - Functions
  - Function Modifiers
  - Events
- Types: integer, string, array, mapping...

We will talk more about this  
in the tutorial session 😊

# Simple smart contract: Lottery

## Contract Lottery (pseudocode)

Init:

Tend := 28 February 2018, \$ticket := 1, pool := {}, pot := 0

Function TicketPurchase():

On receiving \$amt from some party **P**:

Assert \$amt = \$ticket, balance[**P**] ≥ \$amt

balance[**P**] := balance[**P**] - \$ticket

pot := pot + \$ticket

pool := pool ∪ **P**

Contract timer: T

If T > Tend then:

$W \xleftarrow{R} \text{pool}$

balance[**W**] := balance[**W**] + pot

Set up contract parameters and data structures

## Lottery

The lottery purchase function on the contract.

Triggered when some party **P** sends a corresponding transaction on the blockchain

Each contract has a *timer* that is either based on the block index or the real-world time

When certain condition is satisfied, this part automatically executes

# Solidity

- A *contract-oriented*, high-level language for implementing smart contracts
- Targets the Ethereum Virtual Machine (EVM)
- Syntax at <https://solidity.readthedocs.io/>

```
pragma solidity ^0.4.0;

contract SimpleStorage {
    uint storedData;

    function set(uint x) public {
        storedData = x;
    }

    function get() public constant returns (uint) {
        return storedData;
    }
}
```

# Browser-solidity

- An useful integrated development environment (IDE) with solidity
  - Design and run smart contract on a Java VM environment
  - Debugging
- Easier for beginners
  - Will not consume “real” money
  - Need not to set up your local EVM
    - E.g., ethereumjs-testrpc <https://github.com/ethereumjs/testrpc>
- <https://remix.ethereum.org/>

# Related References

- Andrew Miller and Ari Juels and Elaine Shi and Bryan Parno and Jonathan Katz. “Permacoin: Repurposing Bitcoin Work for Data Preservation”, in Proc. of IEEE S&P 2014.
- F. Zhang, I. Eyal, R. Escriva, A. Juels and R. V. Renesse. “REM: Resource-Efficient Mining for Blockchains”, in Proc. of USENIX Security, 2017.
- Aggelos Kiayias and Alexander Russell and Bernardo David and Roman Oliynykov. “Ouroboros: A Provably Secure Proof-of-Stake Blockchain Protocol”, in Proc. of CRYPTO, 2017.
- Yossi Gilad and Rotem Hemo and Silvio Micali and Georgios Vlachos and Nickolai Zeldovich. “Algorand: Scaling Byzantine Agreements for Cryptocurrencies”, in Proc. of SOSP, 2017.
- Eleftherios Kokoris-Kogias and Philipp Jovanovic and Linus Gasser and Nicolas Gailly and Ewa Syta and Bryan Ford. “OmniLedger: A Secure, Scale-Out, Decentralized Ledger via Sharding”, in Proc. of IEEE S&P, 2018.
- Mustafa Al-Bassam and Alberto Sonnino and Shehar Bano and Dave Hrycyszyn and George Danezis. “Chainspace: A Sharded Smart Contracts Platform”. in Proc. of NDSS, 2018.
- Loi Luu and Viswesh Narayanan and Chaodong Zheng and Kunal Baweja and Seth Gilbert and Prateek Saxena. “A Secure Sharding Protocol For Open Blockchains”. in Proc. of ACM CCS 2016.
- Mahdi Zamani and Mahnush Movahedi and Mariana Raykova. “RapidChain: A Fast Blockchain Protocol via Full Sharding”. in Proc. of ACM CCS 2018.